

The Mask and the Mirror

Evaluating the Transition of Character-Based AI from Entertainment to Epistemic Infrastructure

Jacob E. Thomas¹ · Rick Halpern^{1,2} · Shawn Martin^{1,3}

1 The COMPANION Working Group 2 University of Toronto 3 Washington College

Correspondence: JEThomasPhD@gmail.com

Preprint · August 2026 · **Project Repository:** https://github.com/jethomasphd/THE_COMPANION_DOSSIER

Abstract


Role-Playing Language Agents (RPLAs) are transitioning rapidly from entertainment novelties to proposed infrastructure for pedagogy, historical inquiry, and civic deliberation: a transition that demands evaluation standards the field does not yet possess. This article argues that the dominant evaluation paradigm, which measures linguistic fluency, personality consistency, and stylistic naturalness, cannot detect the most consequential failure mode in historically situated role-playing: the collapse of historical figures into anachronistic composites shaped by cultural salience rather than documentary evidence. Historians have long recognized cognate phenomena as presentism and retrospective narrative smoothing. We present three interrelated contributions toward addressing that gap.

First, the *Mask and Mirror* framework distinguishes two conceptions of successful role-playing: the *Mask*, which optimizes for convincingness, and the *Mirror*, which optimizes for documentary fidelity. We argue that current RPLA development has progressively matured the Mask through three paradigm stages (template-based systems, style imitation, cognitive simulation), and that the field now requires a fourth: *epistemic simulation*, in which outputs are bounded by a specified corpus, anchored to a specified temporal moment, and evaluated by domain experts against the evidentiary record. Second, the *Miranda Hypothesis* names the specific failure mode that makes this transition necessary. It holds that the volume and recency of culturally dominant representations in training corpora systematically overwhelm the primary documentary record, producing personas that feel authentic to modern users precisely because they have been contaminated by contemporary moral logic. The hypothesis takes its name from Lin-Manuel Miranda’s *Hamilton*, a culturally dominant representation that has functionally overwritten the documentary Alexander Hamilton in the public imagination.

Third, the *Prism Experiment* is a pre-registered interdisciplinary protocol designed to test the Miranda Hypothesis directly. Developed through collaboration between a data scientist and a domain historian, the protocol deploys Abraham Lincoln across four documented temporal moments, three seeding conditions, and five diagnostic questions crafted to expose anachronistic reasoning, with outputs evaluated against a rubric that prioritizes anachronism detection. Pre-registration of the protocol in advance of data collection is itself a methodological contribution: it establishes a replicable standard for evaluating epistemic fidelity in historically situated role-playing that other researchers can adopt, extend, and challenge. Throughout, we argue that the maturation of persona-simulation technology requires disciplinary expertise (e.g., historical, archival, philosophical, literary) to be recognized as a technical requirement, rather than an optional enrichment.

Keywords: role-playing language agents · historical personas · LLM evaluation · presentism · cultural memory · archival diplomatics · pre-registration · epistemic simulation

Visual Abstract:



1 Introduction

1.1 A Hospital Room

One of us (JET) came to this work not in a research laboratory but in a hospital room. The aim was modest and human: to use a large language model to help a person with advanced dementia speak coherently again. The attempt could not succeed. The documentary record required to anchor such an encounter had not been gathered, and the person for whom the system was being built was, by then, already largely beyond the reach of language. What the model produced, absent anchor material, was not the person’s voice. It was a culturally shaped composite of a “mother” that resembled the person closely enough to mark, with painful clarity, both the promise and limitations of language models.

That failure is the seed of this paper. The methodology was transmuted. The figures whose records permit the construction of legitimate encounters became the practice (i.e., Washington, Hamilton, Jefferson, Franklin, and ultimately Lincoln), while the personal use case became the horizon: it may yet be possible, with adequate archival anchoring, to construct an encounter with the documentary record of a deceased loved one that is grounded rather than fabricated. Everything that follows is, in the end, an attempt to specify how much anchoring is required before that horizon becomes ethically traversable, and to build the instruments that would let us know whether we had arrived. We return to this threshold, and to what it demands of a framework, in Section 2.8.

1.2 The Inflection Point

The current landscape of Role-Playing Language Agents reveals a technology at an inflection point: not between failure and success, but between two fundamentally different conceptions of what success means. A comprehensive survey by Chen et al. (2024) reviews a rapid expansion of persona types (e.g., group-based demographic archetypes, specific historical and fictional characters, individualized personas derived from user data) and documents their deployment across entertainment, education, therapy, and persuasion. A subsequent review by Wang et al. (2026) traces the field’s technical evolution across three paradigm stages: from rule-based template systems, through style imitation enabled by large language models, to what they term *cognitive simulation* (i.e., systems that model personality through psychological frameworks, maintain character state through structured memory), and generate behavior constrained by motivation–situation chains. This third stage represents genuine progress. Systems such as CoSER (Wang et al., 2025) construct motivation-driven behavioral agents from a literary corpus of 17,966 characters drawn from 771 books; PsyMem (Cheng et al., 2025) models characters through 26 quantitative psychological indicators anchored by knowledge-graph memory alignment; and the LIFECHOICE dataset (Xu et al., 2024), comprising 1,401 character decision points from 395 novels, finds that without explicit personality–memory prompting, even frontier models generate markedly inconsistent behavioral choices.

Yet Wang et al. (2026) document a revealing asymmetry at the heart of this progress. The automated evaluation systems now standard in RPLA assessment (e.g., LLM-as-judge paradigms increasingly adopted for scalability) systematically privilege linguistic fluency and stylistic naturalness over fidelity to character settings, and they still lack comprehensive quantitative assessments of personality fidelity and behavioral logic (Wang et al., 2026; Chen et al., 2024). The InCharacter benchmark (Wang et al., 2024), which evaluates personality fidelity through psychological interviews rather than the self-report

scales it critiques, reports state-of-the-art RPLAs achieving up to 80.7% alignment with the human-perceived personalities of target characters. This is an impressive engineering result. But it measures a particular thing, and not the thing this paper is concerned with: it establishes that a model can reproduce (for example) Alexander Hamilton’s Big Five profile, not that it can constrain Alexander Hamilton within his documentary record at a specific moment in his intellectual development. The benchmark answers “does this resemble the character?”; it is silent on “could this character have reasoned this way, at this time, given what the record permits?” That silence is the gap this paper addresses.

These findings reveal a structural orientation in the field that we call the *Mask*. By the Mask we mean the conception of successful role-playing as the production of outputs that feel like a character (i.e., fluent in register, consistent with a personality profile, emotionally responsive in style) without requiring that those outputs be disciplined by the character’s actual documentary record at a specific moment in time. The Mask optimizes for convincingness in the sense of narrative plausibility and coherence familiar from historiographical traditions. It asks: does this sound like the person? It does not ask: is this what the person could have known, believed, or argued at this point in their life? The field has built its measurement apparatus (i.e., benchmarks, reward models, LLM-as-judge protocols) around the Mask, and a Mask can be made arbitrarily convincing without becoming any more faithful.

What is needed, and what this paper proposes, is the Mask’s complement: the *Mirror*. A Mask is worn to be seen; it presents a surface to an audience and is judged by how persuasive that surface is. A Mirror is consulted to see truly; it is judged not by how it looks but by the fidelity of what it reflects back. Where the Mask asks whether a persona is convincing, the Mirror asks whether a persona reasons within the boundaries of what its historical counterpart’s record permits. The constraint is not personality coherence but documentary fidelity; the measure of success is not how convincingly a persona sounds but how faithfully it reasons from a bounded record at a bounded moment. The distinction this paper turns on is therefore simple to state and consequential to operationalize: when a language model speaks *as* Alexander Hamilton, or Abraham Lincoln, or Martin Luther King Jr., is it speaking *from* that figure’s documentary record (Mirrored), or merely *about* it (Masked)?

This distinction is largely immaterial for entertainment. It becomes consequential the moment RPLAs are proposed as infrastructure for domains where epistemic fidelity is not optional. Pedagogical simulations that purport to teach how historical figures reasoned cannot merely produce plausible dialogue; they must produce dialogue constrained by what the documentary record permits, shaped by what the figure had experienced at that moment, and free from frameworks the figure could not yet have possessed. Civic-deliberation tools that convene historical voices to address contemporary crises cannot optimize for engagement without optimizing away the very tension that makes the deliberation valuable. A Hamilton who agrees with a progressive user’s priors is not Hamilton deliberating; it is the model wearing a costume. The evaluation infrastructure Wang et al. (2026) document, optimized for naturalness and personality-scale correspondence, would score that costumed Hamilton highly.

This paper therefore takes a position. **If the dominant failure mode of persona-simulation systems is anachronistic compositing, and the evaluation infrastructure measures personality consistency and linguistic fluency, then the evaluation infrastructure cannot detect the dominant failure mode.** This is not a critique answerable by better automated metrics, more refined LLM-judges, or larger character datasets, because the failure is not in the calibration of those instruments, but in what they

are instruments for. A persona-simulation system without domain expertise in its evaluation loop is a thermometer that cannot read temperature: it returns confident measurements of something else. The maturation of character-based AI from entertainment product to epistemic instrument therefore requires the disciplines trained to read the evidentiary record — history first among them, but also archival science, philosophy, literary criticism, and the relevant social sciences — to be recognized not as optional enrichments but as technical requirements in the design, evaluation, and governance of persona-simulation systems. The relevant expert is not adjacent to this paradigm. The expert is the missing instrument.

Two clarifications guard this claim against overstatement. First, the Mask vs Mirror contrast is an analytic distinction between orientations, not a metaphysical partition between honest and dishonest systems; every act of historical reconstruction, human or machine, is mediated and interpretive, and a Mirror is not a window onto the past but a disciplined reflection of a bounded record. Second, to invoke the domain expert is not to posit a single authoritative evaluator who pronounces the truth. Historical interpretation is irreducibly perspectival and internally plural, and the framework we develop treats that pluralism as a feature: transparency about an evaluator’s commitments is preferable to the illusion of a view from nowhere. The expert is the missing instrument in the sense that some trained reader of the record must be in the loop, not in the sense that any one reading is final.

1.3 A Note on Terminology

Because this article addresses readers in machine learning, history, and the information sciences at once, and because those communities use several of the same words to mean different things, we fix a small working vocabulary before proceeding. We mean these definitions to be serviceable rather than canonical; specialists in each field will recognize that fuller treatments exist.

Role-Playing Language Agent (RPLA). The field’s standard term for a system that produces in-character dialogue as a persona. We retain it when referring to the literature.

Role-Playing Language System (RPLS). Our own unit of analysis (Section 2.1): the entire configured encounter (i.e., prompt, anchor material, temporal specification, model, and human interlocutor) rather than the model alone.

Corpus. Here, a bounded set of primary or secondary documents deliberately placed in the encounter; in addition to, as in some computational usage, the model’s entire training distribution.

Anchoring. The deliberate constraint of a persona to a moment (**temporal anchoring**) and to a document set (**corpus boundedness**).

Digitization. We use this in the computational sense of converting content into machine-readable text at scale; archivists reserve the word for the making of faithful digital facsimiles of specific holdings, and we flag the divergence where it matters.

Accessibility. We use this to mean the breadth of people and institutions able to *build* an RPLS without specialized compute or expertise (Section 2.6); it is distinct from the disability-access sense the word carries in library and information practice.

Provenance / respect des fonds. Archival principles requiring that records be maintained according to their origin and original order, preserving the chain of custody and context that authenticates a document.

Diplomatics. The information-science discipline concerned with the form, genesis, and authenticity of documents, which we draw on in Section 2.5.

Four organizing metaphors recur throughout, and we name them together here so the reader can hold them in view. The **Mask** is the orientation that optimizes convincingness; the **Mirror** is the orientation that optimizes documentary fidelity. **Miranda distortion** is the specific failure the Mask produces (i.e., a culturally salient composite displacing the documentary figure) and the **Miranda Hypothesis** is the empirical claim that current systems produce it by default. The **Prism** is the experimental instrument that attempts to refract a single composite figure into the distinct, temporally bounded figures the record actually contains. Mask and Mirror are conceptions of success; Miranda is a failure mode; the Prism is an instrument. We have tried, in revising this article, to deploy this vocabulary where it does analytic work and to resist it where plainer language serves better.

1.4 Three Contributions, and Their Evidentiary Status

We wish to be exact, at the outset, about what we have demonstrated and what we propose, because the distinction governs how the rest of the paper should be read. The Mask/Mirror framework and the Miranda Hypothesis are arguments: a normative account of what role-playing should be for, and an empirical claim, with a specified mechanism, about what unconstrained systems produce. The Prism Experiment is a pre-registered research design, not a completed study; its diagnostic questions, corpus pairings, rubric, and the historian’s sealed vignettes are presented here as a standard established before the data exist. We frame pre-registration not as a hedge but as the methodological strength it plainly is, and having established that status once, here, we proceed assertively rather than re-qualifying it at every turn.

This paper makes three contributions. First, the *Mask and Mirror framework* formalizes the distinction introduced above and identifies the architectural and evaluative shifts required to move from one orientation to the other. The three-stage evolution Wang et al. (2026) describe — templates, style imitation, cognitive simulation — is, in our framing, the progressive sophistication of the Mask. Each stage produces more convincing characters; but convincingness and fidelity are independent properties, and a system can score perfectly on personality consistency, behavioral coherence, and memory persistence while producing a figure who reasons from knowledge their historical counterpart never possessed. The Mirror requires a fourth stage we term *epistemic simulation*, in which the governing constraint is not personality coherence but documentary fidelity: outputs bounded by a specified corpus, anchored to a specified temporal moment, and evaluated by domain experts against the evidentiary record rather than by automated scorers against personality scales.

Second, the *Miranda Hypothesis* names the specific failure mode that makes the transition necessary. The hypothesis takes its name from Lin-Manuel Miranda, whose *Hamilton* (2015) constitutes its paradigm case: a culturally dominant representation so saturating, so rhetorically powerful, and so morally legible to contemporary audiences that it has functionally overwritten the documentary Alexander Hamilton in public imagination and, we argue, in the training corpora of every frontier language model. We hypothesize that without temporal anchoring and corpus constraint, language models default to culturally salient composites: figures who speak in a historically plausible register but reason from an undifferentiated amalgam of everything the corpus contains about them, including interpretive frameworks, retrospective judgments, and cultural mythology that postdate the figure’s own lifetime. The composite Hamilton knows

he will be the subject of a Broadway musical; the composite Lincoln has already read the Gettysburg Address before he has written it. The hypothesis predicts that this compositing is not a bug to be patched but a default mode emerging from the architecture of autoregressive language modeling itself, and that overcoming it requires external constraint structures the current RPLA paradigm does not provide.

Third, the *Prism Experiment* is a pre-registered interdisciplinary protocol designed to test the Miranda Hypothesis directly. Developed through collaboration between a behavioral epidemiologist and data scientist (JET) and a historian of American social and political history (RH), with archival and information-science framing contributed by a librarian and academic theologian (SM), the experiment tests whether the temporal anchoring of historical personas can discipline language-model outputs from composite toward documentary fidelity. The collaboration proceeded through four phases: a published open-source protocol and proof-of-concept deliberation (Thomas, 2025); a domain expert’s critical review producing five interventions (Thomas & Halpern, 2026); and the historian’s formal contributions of temporal moments, corpus pairings, diagnostic questions, a priori vignettes, and an evaluation rubric (Halpern, 2026). We publish the protocol before reporting results: the design, the corpus pairings, the diagnostic questions, the vignettes (held under cryptographic seal, by the mechanism specified in Section 4.4), and the weighted rubric are presented here as a complete pre-registration, to establish the evidentiary standard before the data exist, to invite scrutiny and replication from both the AI and the humanities communities, and to model an interdisciplinary transparency the RPLA field (for all its technical sophistication) currently lacks.

2 The Mask and the Mirror

We turn from diagnosis to formalization. The previous section named two orientations toward role-playing language agents (i.e., the Mask, which optimizes for convincingness, and the Mirror, which optimizes for documentary fidelity) and argued that the field’s evaluation infrastructure is constituted around the former. Here we develop the framework into a working taxonomy, articulate its technical commitments, and identify the two architectures by which RPLAs are presently constructed, only one of which can in principle yield a Mirror (See Figure 1).




Figure 1. The Mask and the Mirror framework. Two orientations toward role-playing language systems: the *Mask*, which optimizes for convincingness, and the *Mirror*, which optimizes for documentary fidelity. The figure also situates the field’s development across four paradigm stages, culminating in epistemic simulation.

2.1 The Role-Playing Language System as the Unit of Analysis

The literature speaks of role-playing language agents. We adopt instead the term Role-Playing Language System (RPLS) to mark a deliberate shift in the unit of analysis: from the *agent* (an autonomous synthetic entity) to the *system* (the entire structured encounter that includes its human, its prompt, its anchor material, and its model). The agent-framing presumes the persona is a property of the model itself; the system-framing locates the persona in the configuration that brings model, document, and interlocutor into

relation. This is not an unfamiliar move. It echoes a methodological reorientation long since accomplished in the human sciences — most influentially in the *Annales* tradition (Bloch, 1953; Braudel, 1972), which displaced the individual as the primary unit of historical analysis in favor of structures, relations, and the long-run conditions that shape action — and it resonates with the sociology of knowledge, which locates the production of facts not in solitary minds but in institutions, instruments, and circuits of citation and exchange (Merton, 1973). To relocate the persona from the model to the configuration is, in this sense, a familiar reorientation rather than an exotic one.

A minimal RPLS consists of five components: a structured prompt that supplies framing and constraints; anchor material drawn from primary documents that constrain and authenticate the persona; a temporal anchor specifying the moment in the figure’s life from which the encounter is to speak; an off-the-shelf large language model that reads and speaks through the materials (the voice, but not the mind, of the encounter); and a human interlocutor who curates what enters the encounter, retains interpretive custody over what emerges from it, and brings relational and contextual knowledge as a check on the model’s drift. Defined this way, an RPLS is not a simulation of a person; it is a configured encounter with a person’s record, at a specified moment in the life that record describes.

2.2 The Temporal Anchor

The temporal anchor is not an embellishment of this definition but a foundational commitment of it, even as temporal anchoring itself remains mediated by corpus selection and interpretive framing. The persona an RPLS instantiates is not a person but a person-at-a-moment. Alexander Hamilton in 1789, in the months following ratification and on the eve of his appointment to the Treasury, is a different epistemic object than Alexander Hamilton in 1804, riding out to Weehawken with twenty-five years of accumulated experience, defeat, and grievance. To collapse them is to produce neither in any strictly historical sense. What is produced in their place is a composite persona whose voice belongs to no moment in the figure’s life but to the entire arc of the figure’s representation in the training corpus, undifferentiated by time. This collapse is the technical mechanism of Miranda distortion, developed in Section 3; it is also the failure the temporal anchor is built to foreclose. An RPLS without a temporal anchor cannot in principle yield a Mirror, because what it instantiates is not a moment of reasoning but a tableau of legacy.

The distinction between simulation and encounter follows directly. A simulation purports to instantiate the person; an encounter purports only to make the record speak in dialogue, at the moment the historian, the curator, or the grandchild has chosen to convene it. The first claim is strong and, we will argue, ill-formed; the second is rhetorical and tractable. The Mask is the orientation that mistakes the second for the first.

2.3 Two Architectures

Among present-day approaches to constructing an RPLS, two architectures dominate. The first places anchor material directly into the model’s context window at inference time, alongside the structured prompt, drawing on the technical lineage of retrieval-augmented generation (Lewis et al., 2020). The second uses anchor material as training data, adjusting the model’s weights through supervised fine-tuning, parameter-efficient adaptation, or continued pretraining — the approach exemplified by Character-LLM (Shao et al., 2023) and developed at scale in CoSER (Wang et al., 2025). Real systems often blend the two

(retrieval-augmented inference layered over fine-tuned bases, or tool-augmented retrieval at runtime), but hybrid architectures inherit the liabilities of fine-tuning without escaping them, as we explain below.

On first inspection the two architectures appear to be alternative implementations of the same goal. We argue they are not: they embody opposed conceptions of what role-playing is, and they yield epistemically different objects. In the **context-window architecture**, anchor documents enter the model's context at the moment of encounter and leave it at the encounter's end. The model reasons over the documents in real time; the persona is interacted with, through, and against the anchor material. Documents remain documents (i.e., they are not transformed, and they survive the session intact). Provenance is preserved end to end: at any moment one can ask which documents are in the window. The encounter is reversible: the session ends, the documents persist, the next encounter begins fresh. The human retains interpretive custody, curating what enters the window, interrogating the output, and bringing relational knowledge as a check on drift.

In the **fine-tuning architecture**, anchor documents are consumed as training data. The model's weights are adjusted through exposure to the corpus, and the persona becomes an aspect of the model's parameter space: distributed, indistinguishable from everything else the model has learned, inseparable from the base model's priors. Documents do not survive the encounter in any meaningful sense; they are dissolved into weights whose relationship to the original corpus is statistical and unverifiable from outside. The human loses interpretive custody: gradient descent, not the curator, determines which features of the corpus are retained, which are smoothed away, and which are amplified by interaction with priors. The encounter is irreversible in the meaningful sense, because weights cannot be unbaked. The first architecture treats role-playing as performative; the second treats it as ontological. Fine-tuning attempts to make the model be the persona; the context window allows the model to speak through the persona's record. These are not implementations of the same goal but answers to different questions about where a persona lives.

2.4 Where Personas Live: Performance, Provenance, and a Productive Category Mistake

To say that role-playing is performative is to invoke a tradition with substantial theoretical lineage. From Goffman's analysis of self-presentation as situated performance (Goffman, 1959), through Searle's speech-act theory (Searle, 1969), to Butler's account of identity as iterated performance (Butler, 1990), role is not a property an entity has but an act it performs — in a scene, before an interlocutor, under constraints. Role is not in the actor; it is in the configuration that includes actor, script, stage, and audience. A persona, by this logic, can no more be located in a model's weights than Hamlet can be located in Laurence Olivier's body. The persona is an event that occurs when model, prompt, anchor material, and interlocutor convene. The point has a longer pedigree still: Foucault (1977) reframed the author not as a stable origin of meaning but as a *function* produced within discourse; Steedman (2001) recast the archive as a site of encounter shaped by desire and interpretation rather than passive preservation; and Portelli (1991) showed how meaning — especially in oral sources — emerges through performance, memory, and narration rather than residing in fixed content. Across these traditions, what a figure means is an effect of the encounter, not a possession of the document.

The fine-tuning architecture, by attempting to internalize the persona into the model's parameters, commits what we will call a category mistake: it tries to relocate something that lives in the configuration into something that lives in the model. We use the phrase advisedly, and with one qualification historians

will rightly press. Productive category mistakes are common and often analytically useful: scholars routinely speak of “the voice” of an author while knowing that voice is an effect of genre, mediation, and survival. So the relevant question is not whether fine-tuning is simply mistaken, but whether it is a cruder way of stabilizing something interpreters also try to stabilize at the cost of the very thing that makes such stabilization legitimate. That cost is auditability. The decisive difference between the two architectures is not distortion as such (all reconstruction distorts) but diagnosability: in a context-window configuration, drift toward composite can be traced to a source and corrected; in a fine-tuned model, it is submerged into parameters and can no longer be said why something appears as it does. This maps almost exactly onto the archival concern with provenance and transparency, to which we turn next: interpreters tolerate interpretation so long as the evidentiary chain remains visible. What is unsettling about fine-tuning a persona is not that it interprets but that it severs the chain.

This has measurable consequences for the Miranda Hypothesis. When persona is performative, Miranda distortion is managed by configuration: anchor material in the context window disciplines drift toward composite, and provenance allows the source of any drift to be diagnosed. When persona is ontological, distortion is amplified: a thin personal signal is layered over the vast cultural sediment already present in the base model’s weights, and the signals interact in ways that no longer admit external audit. The user perceives a more convincing Hamilton precisely because the system has merged Hamilton’s documentary record more deeply with everything else the corpus contains about him — including, and especially, *Hamilton* the musical. Fine-tuning suppresses Miranda distortion at the surface while amplifying it beneath. The context-window architecture does not eliminate the base model’s priors; it makes the source of any drift diagnosable. The Mirror is not a system without priors. It is a system in which priors can be interrogated rather than absorbed.

None of this is unprecedented; it is newly intensified. There are older moments when cultural representations began to overwrite archival ones — the nineteenth-century historical novel, or Shakespeare’s histories shaping the perception of English kings. What is new is that fine-tuning collapses those layers internally, making representation and record harder to disentangle than any prior medium allowed. We do not claim that fine-tuning has no role in RPLS construction. It is appropriate for the structural features of an era — period orthography, document-genre conventions, citation forms — that is, for the grammar of a context rather than the voice of a person. Historians and information scientists are generally more comfortable abstracting structures (genre, discourse convention, orthography) than persons, precisely because the former are already understood as distributed and impersonal. Fine-tuning works where we already accept generalization, and fails where we insist on irreducible specificity. The argument here concerns persona, not period.

2.5 Archival Ethics and Diplomats: Documents as Sites of Return

The argument to this point has been conceptual and technical. Its full significance emerges only when placed in conversation with archival theory — and, crucially, with archival practice, where the stakes are not merely methodological but ethical, and are negotiated in the everyday handling of historical materials. We frame this section through *diplomats*, the information-science discipline concerned with the form, genesis, and authenticity of documents, because diplomats supplies exactly the analytic vocabulary the comparison requires: it treats a document not as a free-floating string of content but as an act situated within the context of its creation, custody, and use (Foscarini, 2012; Blouin, 1996). That

orientation has been carried into the digital environment by a substantial literature on the authenticity and trustworthiness of records held and transformed by computational systems (Duranti & Rogers, 2019; Jansen, 2015; Suderman et al., 2025), which is the literature most directly relevant to what fine-tuning does to a corpus.

The governing intuition is old. In *Archive Fever*, Derrida (1996) characterized the archive not as a passive repository but as a site of return — a formulation echoed in scholarship that understands archives as sites of epistemic production and administrative ordering rather than neutral storage (Stoler, 2009), and as locations whose value lies in the enduring, accumulative life of documents within them (Steedman, 2001). The document is meant to be visited. It survives every encounter and remains available to the next reader, the next question, the next generation. The archive's value is precisely that it is *not consumed by use*: what was read can be re-read; what one reader excluded the next can include.

This principle is not abstract; it is built into the standard protocols of archival work. In a reading room a researcher is given a single box at a time; documents must be handled carefully, kept in order, and returned to their folders exactly as found; gloves may be required, pencils replace pens, photography is regulated. These practices are not only about physical preservation. They encode an ethical commitment — that the document, once encountered, must remain available, materially and intellectually, for future encounters — and they instantiate the principle archival theory calls *respect des fonds*, that records be maintained according to their origin and original order (Duchemin, 1983). Diplomatics sharpens the point: what authenticates a document is not its content alone but the context of its genesis — who created it, for whom, when, in what administrative or relational circumstance — and that context can itself shift over time and must be preserved alongside the text (Blouin, 1996). A letter in a family collection is not simply a text; it is a node in a network of relations (i.e., between correspondents, across time, within a life) and the archivist's task is to make that network legible without collapsing it. In this sense the metadata of a document (i.e., its addressee, its date, its place in a sequence) is as much a part of the record as the words it carries.

From this vantage the two RPLS architectures stand in sharply different relation to archival ethics. The context-window architecture aligns closely with these principles: documents enter the interpretive encounter as discrete entities — analogous to a file requested from a box — and leave it unchanged, retaining their integrity, provenance, and availability for future engagement. One researcher's question does not preclude another's; one interpretive frame does not overwrite the document's structure; what is excluded in one reading can be included in the next. The document persists as a site of return. The fine-tuning architecture, by contrast, operates according to a logic with no clear analogue in archival practice. It resembles not the consultation of records but their irreversible transformation: to fine-tune on a corpus of letters or diaries is, in effect, to disassemble them — stripping away order, context, and relational structure — and recombine their contents into statistical parameters from which no original letter can be requested, traced to its archival location, or situated among its neighbors. The chain of provenance is broken. In archival terms it is akin to taking a box of correspondence, shredding the documents, and using the fragments to produce a composite from which no original could be reconstructed or revisited — a violation not only of preservation but of original order and the possibility of future reinterpretation. Derrida's own most apposite contribution here is not his theory of the archive in general, about which archivists reasonably differ, but his account of archival violence: the destruction of records and memory, often performed in the

name of efficiency or order, that erases the very thing the record was created to hold. Rendering a curated archive into model weights that dissolve its historical context is, in this precise sense, a candidate form of archival violence, accomplished not by malice but by extraction.

The charge intensifies for personal and familial materials such as letters, journals, recorded voices, or other intimate residue of a life. Archives process such materials with particular care: donor agreements specify conditions of access; archivists consult families about sensitive content; oral-history collections impose embargoes and consent restrictions; personal papers may be partially closed or redacted. These practices recognize that such artifacts were created as acts of communication and self-expression, not as raw material for transformation. To treat them instead as training data is to change the nature of how they were intended to be used, and so to raise a genuine ethical question: should a letter to a family member be handled as a communicative artifact, embedded in provenance and access conditions, or as an instrumental input — a quarry of historical vocabulary? To convert the former into the latter bypasses the very structures — provenance, access conditions, contextual framing — that define the artifact’s ethical standing within the archive. The contemporary diplomatics literature on records in the cloud frames exactly this problem as one of preserving contextual integrity when documents pass into computational systems that would otherwise sever it (Duranti & Rogers, 2019; Jansen, 2015; Suderman et al., 2025).

This is the connection to the paper’s central argument, and it is worth making explicit. The Mirror asks that a historical person be granted integrity appropriate to a specific moment in time: that the figure not be made to reason from premises he could not yet have held. Diplomatics asks the cognate thing of a document: that it be preserved within, and interpreted through, the historical context of its creation. The two commitments are the same commitment, applied to person and to record: fidelity to a thing as it actually stood at a moment, against the pressure to flatten it into a timeless, frictionless, present-tense composite. An architecture that preserves documents as sites of return is therefore not merely the more cautious engineering choice. It is the choice that honors, at the level of the record, exactly the discipline the Mirror demands at the level of the persona. The context-window architecture treats primary documents as things to be visited; the fine-tuning architecture treats them as things to be consumed. The distinction is not aesthetic. It is ethical, and it is methodological, and the two are not separable.

We state this strongly while conceding what a careful archivist would insist upon: archival processes are complex, and these contrasts simplify them. Not every fine-tuning use is a desecration, and not every context-window use is virtuous; the relevant ethics are matters of consent, scale, and purpose as much as of architecture. Our claim is narrower and, we think, defensible: *as a default orientation toward the documentary record*, the context-window architecture preserves what archival theory and diplomatics are structured to protect, and the fine-tuning architecture is structured to dissolve it.

2.6 Accessibility as Methodology

The two architectures differ along a further axis that the technical literature consistently treats as secondary but that the humanities has long argued belongs at the center of any framework: accessibility. We use the term here in a specific sense — the breadth of people and institutions able to build an RPLS at all — and we flag at the outset that this is not the sense the word ordinarily carries in library and information practice, where accessibility concerns access for people with disabilities. Both senses matter; only the first is our subject in this subsection.

Fine-tuning requires compute infrastructure (GPU access, training pipelines), technical expertise (dataset curation, hyperparameter selection, evaluation harnesses), institutional access (most individuals and most humanists cannot fine-tune frontier models), and corpus scale (small personal archives produce overfitting artifacts rather than authentic personas). It is, in short, an institutional capability. The context-window architecture, by contrast, requires literacy, a set of documents, and access to any modern frontier language model, including free-tier services. It is a kitchen-table capability.

This is not a matter of operational convenience. It determines which forms of inquiry the technology can support and, most critically, which communities can author it. A doctoral student in early American history; a community archivist documenting a regional figure; a secondary-school teacher building a classroom encounter with a abolitionist pamphlet; a museum educator preparing a gallery interpretation; a genealogist working a family collection; an Indigenous community convening the recorded words of an elder under its own protocols of access; a grandchild sitting with a grandmother’s letters... these are precisely the users for whom RPLS methodology has the most to offer, and precisely the users for whom fine-tuning is structurally inaccessible. An RPLS that can only be constructed by computational laboratories is not infrastructure for the humanities; it is infrastructure for whoever owns the laboratory.

The framework’s commitment to accessibility is therefore not a populist gesture appended to a technical argument. It is itself a technical argument. The architecture that admits the most diverse population of curators is the architecture most likely, over time, to surface the documentary anchorings the field most needs, and to be governed by the epistemic standards of the disciplines whose materials are placed at its center.

2.7 From Mask to Mirror: Four Paradigm Stages

Wang et al. (2026) trace the development of role-playing language agents through three paradigm stages: rule-based templates, style imitation, and cognitive simulation. We accept this taxonomy and propose a re-reading of it. The three stages are not a sequence of failed approximations followed by a recent success. They are a sequence of progressively more sophisticated Masks, each more convincing than the last, each evaluated by instruments calibrated for convincingness. Cognitive simulation is the most sophisticated Mask the field has built. It is not, by virtue of its sophistication, a Mirror. The stages are best read through what each takes as its primary constraint.

In the template stage, the constraint is structural: the system reproduces a fixed pattern of dialogue (i.e., canned responses keyed to user inputs) and the persona is a thin lacquer over a finite-state machine. Convincingness is low because the persona’s fixity is exposed by any unanticipated turn. In the style-imitation stage, the constraint is registral: the system reproduces the surface features of a figure’s voice (i.e., vocabulary, cadence, characteristic tic) without modeling what the voice would say in situations the corpus does not contain. Convincingness rises because the voice generalizes; fidelity does not, because the generalization is statistical mimicry of style rather than reasoning from substance.

In the cognitive-simulation stage, the constraint is psychological: the system models personality through trait inventories, maintains state through memory architectures, and generates behavior through motivation–situation chains (Wang et al., 2024; Cheng et al., 2025; Xu et al., 2024). Convincingness rises again because the persona now has a structured inner life. But the constraint remains internal to the model. It is the shape of the persona’s mind, not the discipline of its record. A cognitively simulated Hamilton has

a convincing Big Five profile and a convincing motivational architecture; he does not, by virtue of this construction, possess any constraint preventing him from quoting *The Federalist* in 1785, two years before those papers were written, or from reasoning about his legacy in language drawn from a Broadway musical that postdates his death by two centuries.

We propose a fourth stage. In *epistemic simulation*, the governing constraint is no longer internal to the model at all. It is external, documentary, and temporal. Three commitments distinguish epistemic simulation from its predecessors. **Corpus boundedness**: the persona’s reasoning is licensed only by a specified corpus of primary documents; what the figure could have known, believed, or argued is bounded by what the corpus permits. The model is not a substitute for the archive; it is a reader of the archive. **Temporal anchoring**: the persona is instantiated at a specified moment in the figure’s life, and knowledge, framing, and language postdating that moment are out of bounds regardless of how culturally salient they have since become. The figure does not know what he has not yet learned. **Expert-loop evaluation**: outputs are evaluated against the evidentiary record by domain experts whose training is in the discipline whose materials the persona claims. Automated metrics (i.e., LLM-as-judge protocols, personality-scale correspondence, fluency benchmarks) are at most diagnostic; they cannot adjudicate fidelity, because fidelity is a judgment about the relation between an output and a record, and that judgment requires knowledge of the record.

These commitments are not additive refinements of cognitive simulation. They constitute a different orientation. Cognitive simulation asks: does this system produce outputs consistent with the figure’s mind? Epistemic simulation asks: does this system produce outputs the figure could have produced at this moment in the figure’s life, evaluable against the figure’s record by someone who has read it? The first question can be answered by automated instruments operating on the model alone. The second cannot. It requires the corpus, the temporal anchor, and the expert: in a phrase, the system, not the agent. The Mirror, then, is not a single system but the orientation that makes epistemic simulation legible as a category. To build a Mirror is not to choose better priors, larger context windows, or more refined personality scales. It is to redefine what role-playing is for, and therefore what it means to do it well. The Prism Experiment, treated formally in Section 4, is the first attempt to build evaluative instruments adequate to the redefinition.

2.8 The Stakes of Failure

We close this section by returning to the threshold named in Section 1.1, because it specifies the condition under which this framework counts itself a success or a failure. The hospital-room origin (i.e., the attempt to give language back to a person whom dementia had placed beyond its reach, and the culturally shaped composite the model produced in place of the person) is not biographical color appended to a technical argument. It is the threshold of the framework’s accountability.

An RPLS architecture that produces convincing fabrications when the persona is the user’s own beloved is not a research artifact with technical limitations; it is a violation. Every analytic move in this paper (i.e., the insistence that documents remain documents, that interpretive authority remain with the human, that the encounter remain reversible, that fidelity be measured against a record rather than against fluency) is a constraint motivated by the recognition that the technology must be evaluated at the threshold of its hardest use case, not its easiest. One does not want a model that *is* one’s mother. One wants a model that can speak with one’s mother’s documents in the room. A framework that cannot meet a grandchild at

a grandmother's letters is not a framework at all; it is another failed product. The figures whose records permit legitimate encounters (i.e., Washington, Hamilton, Jefferson, Franklin, and the Lincoln of the experiment that follows) are the practice. The grounded encounter with the documentary record of a loved one remains the horizon, and the experiment of Section 4 is the instrument by which we measure how much anchoring is required before that horizon becomes ethically traversable.

3 The Miranda Hypothesis

The Mask and Mirror framework names a structural orientation. The Miranda Hypothesis names the failure mode that the orientation must overcome. Where the framework asks how role-playing language systems should be conceived, the hypothesis asks what they actually produce when conceived poorly, and by what mechanism. The framework is a normative argument about what counts as success in an RPLS; the hypothesis is an empirical claim about the default output of any RPLS that does not impose external constraint. The two arguments depend on each other. Without the hypothesis, the framework is an aesthetic preference; without the framework, the hypothesis describes a problem with no theory of repair. This section establishes the paradigm case from which the hypothesis takes its name, states the hypothesis in three clauses, demonstrates it at the surface, situates it in the three theoretical traditions it inherits from, engages the principal objections to its mechanism, marks its persistence inside the time-locked architectures that have approached the same problem from the substrate side, and closes by translating it into four falsifiable sub-claims that the Prism Experiment is designed to test.

3.1 The Paradigm Case: How Hamilton Replaced Hamilton

Lin-Manuel Miranda's *Hamilton: An American Musical* opened on Broadway in August 2015. It has since grossed more than a billion dollars, won eleven Tony Awards and the Pulitzer Prize for Drama, and reached an audience expanded by a 2020 Disney+ release into the tens of millions of households; its cast recording was among the highest-debuting and longest-charting in Billboard history, and the Hamilton Education Program has placed the musical in the curriculum of more than two hundred thousand American secondary-school students (Miranda, 2015). Scholars have read it as a dominant popular representation of early American history and vigorously debated its interpretive authority — most prominently in a dedicated collection of historians' essays assessing what the musical got right, what it got wrong, and why it matters (Romano & Potter, 2018), and in influential critiques of its racial casting and archival politics (Monteiro, 2016; Schocket, 2015). By any reasonable metric of cultural penetration, it is among the most widely encountered representations of an American historical figure in the past half-century.

The displacement is documented at the level of the historic site. Within a year of the premiere, the Schuyler Mansion State Historic Site in Albany recorded a near-tripling of annual visitors and a markedly younger demographic (Skala, 2018). These visitors arrived, in the staff's documented experience, already possessing a body of "facts" about the Schuylers and the Hamiltons, many inaccurate, some inversions of the record. They believed the Schuyler family had three children (the musical centers Angelica, Eliza, and Peggy) when in fact the Schuylers had fifteen, of whom eight survived to adulthood (Chernow, 2004); they believed Alexander Hamilton an unambiguous abolitionist, a persistent myth contradicted by the record (Chan, 2004; Serfilippi, 2020); and they assumed the register of his correspondence resembled the lyrical register of contemporary musical theater. The mansion's interpretive staff began the attritional work of explaining that much of what audiences had learned from the musical was not what the documentary record contained (Schuyler Mansion State Historic Site, 2016).

This is not unique to one site. Monteiro (2016), in what has become a canonical critique of the musical's historiography, documented a series of structural distortions: the elision of Hamilton's complicity in slavery, the casting of the Founders as protagonists of a civic mythology in which racial and economic conflict are absorbed into stylized lyrical conflict, and the absorption of the Founders into a contemporary

moral framework in which their dilemmas are legible to twenty-first-century audiences in ways they were not legible to themselves. Schocket (2015) places the musical within a longer arc he terms *Founders Chic*: a decades-long shift in popular American historiography toward biographical celebration in which the figures of the founding period are repurposed as vessels for the biographer's contemporary politics. The musical, on this reading, is not an outlier but the exemplar of a known displacement.

Nor is the pattern confined to the American founding. The phenomenon recurs wherever a culturally dominant retelling saturates the representation of a historical period. Popular screen treatments of the Tudor court have durably reshaped public understanding of sixteenth-century English politics, and the longue-durée mythology of British exceptionalism — the islands' supposed apartness from the Continent — has been shown to operate as a live political resource, mobilized in the rhetoric of the Brexit referendum and traceable to imperial-era self-conception rather than to the documentary particulars of any single reign (Sèbe, 2023). Furthermore, Roman antiquity, refracted through a long line of cinematic and televisual spectacle, reaches contemporary audiences pre-shaped by representations that postdate the period by two millennia. In each case the structure is the same as in the Hamilton case: a salient, morally legible, heavily replicated representation stands between the audience and the record. We develop the Hamilton case in detail because it is the best-documented instance, not because it is the only one; the hypothesis is general, and the experiment of Section 4 deliberately tests it on a different figure, shaped by a different composite, to guard against mistaking the paradigm case for the phenomenon.

What we add to this critique is a single observation about its consequences for the construction of language models. Every frontier model in production at the time of writing has been trained on a web corpus that postdates 2015 and that contains, on any reasonable measure of representational density, more *Hamilton-the-musical* content than *Hamilton-the-documentary-record* content. *The Federalist* exists as a fixed corpus of roughly 175,000 words; the body of post-2015 secondary content occasioned by the musical — reviews, fan analyses, lyric annotations, educational materials, news coverage, social-media discourse, derivative works, and scholarly responses to all of these — exceeds that volume by orders of magnitude and is, on every metric of cultural salience available, more recent and more recurrent. The musical is not merely present in the training corpus; it is dominant in the corpus's distribution of representations of Alexander Hamilton. When a contemporary model produces an utterance about Hamilton, the statistical character of that utterance has been shaped by this dominance. The musical has not only replaced Hamilton in public memory, as Monteiro and Schocket document; it has replaced him in the model's parametric memory, by simple force of representational frequency.

This is the displacement the Miranda Hypothesis names. The hypothesis honors Miranda not as villain — the musical is a substantial work of art operating within a long historiographical tradition it did not invent — but as the figure who supplied the paradigm case: a representation so saturating, so rhetorically powerful, and so morally legible to contemporary audiences that it has functionally overwritten the documentary record in two registers at once. The first is the public-memory overwrite Monteiro and Schocket document. The second is the corpus-memory overwrite, by which language models inherit not the documentary Hamilton but the cultural composite. The first is a humanities phenomenon; the second is a technical one; and the argument of this section is that the second is produced by the first, mediated by the architecture of contemporary language models, and that no method internal to those models can fully repair it.

3.2 Stating the Hypothesis

We state the Miranda Hypothesis in three clauses. The clauses are sequential rather than independent: the input clause specifies a property of the corpus, the mechanism clause specifies the procedure by which the model converts that property into output, and the output clause specifies the character of the resulting persona (See Figure 2).




Figure 2. The Miranda Hypothesis. A schematic of the hypothesis: asymmetry between primary documentary record and culturally dominant secondary representations, transformation of that asymmetry through autoregressive language modeling, and the production of a composite persona shaped by cultural salience rather than documentary fidelity.

First, the input clause. In the training corpora of contemporary frontier language models, the volume and recency of culturally dominant secondary representations of historical figures systematically exceed the volume and recency of those figures’ primary documentary record (Benjamin, 1935/1968; Trouillot, 1995; Dodge et al., 2021). This asymmetry is structural, even where its degree varies across figures and corpora. It is produced by the joint operation of digitization, which favors recent text (Luccioni & Viviano, 2021); platform incentives, which favor culturally salient text (Gillespie, 2014); and aggregation, which favors text linked to and discussed by other text (Brin & Page, 1998). It holds for virtually every culturally salient historical figure who is also the subject of a popular contemporary retelling (Beer, 2013).

Second, the mechanism clause. The generation of historical personas in contemporary language models is a structural product of autoregressive next-token prediction, functioning as parametric reconstruction rather than document retrieval (Brown et al., 2020). During training, a model compresses web-scale corpora into parameters, learning conditional probability distributions that lack any architectural capacity to distinguish primary documentary records from secondary scholarship and tertiary cultural mythology (Bender et al., 2021). To the model, a 1789 letter and a 2019 viral post about a pop-cultural retelling are alike sequences of tokens. Because modern training datasets are saturated with recent, heavily replicated, derivative content, the statistical weight of contemporary cultural narratives overpowers sparse primary sources (Dodge et al., 2021), and the model’s output defaults to a weighted composite that prioritizes algorithmic salience over historical fidelity. We argue below (Section 3.5) that standard post-training alignment does not correct this tendency and, under current paradigms, tends to reinforce it.

Third, the output clause. The resulting persona is a composite: internally consistent in voice, historically plausible in register, and morally legible to contemporary users, yet not corresponding to the historical figure at any verifiable moment in the documented life even though outputs may partially align with documented positions. Because autoregressive models excel at generating fluent text optimized for narrative coherence rather than epistemic grounding (Ji et al., 2023; Weidinger et al., 2022), the model faithfully reconstructs the figure’s cultural afterlife: the aggregate, present-day rendering encoded by the training corpus, standing in for lived historical reality with a synthesized cultural memory (Assmann, 2011). This composite is persuasive precisely because it mimics the dynamics of public memory itself. Historical narratives are subjected to “smoothing” over time, a process by which contradictory facts and abrasive realities are sanded down by repeated retelling into cohesive, consumable archetypes (Lowenthal, 1985). When a model reproduces this smoothed narrative, it aligns with the human cognitive bias toward narrative coherence and familiar tropes over factual complexity (Lin et al., 2022). *Miranda distortion* (i.e., the phenomenon by which a highly salient fictionalized representation eclipses the record) therefore does not register to the user as an error. Because the composite matches the shape of the inherited myth, the distortion feels authentically, incontrovertibly real.

We name the composite *Miranda distortion*, after the paradigm case, and the failure mode that produces it the *Miranda Hypothesis*. The hypothesis is empirical, not metaphorical. It claims that this is what current systems do, by default, in the absence of intervention. It is not the claim that current systems are useless or that they fail at some easier task; it is the claim that they succeed at producing convincingness and fail at producing fidelity, that the success and the failure are produced by the same mechanism, and that any evaluation paradigm that does not distinguish the two will register only the success.

3.3 Demonstration: The Composite at the Surface

Two illustrative observations make the architectural argument concrete. They are not experimental data (that is the work of Section 4) but demonstrations adequate to the theoretical claim and recognizable to any reader who has interrogated a frontier model about a culturally salient historical figure.

The first observation is registral. A frontier model prompted to speak as Alexander Hamilton will produce dialogue whose register draws audibly on the cadence and emotional palette of the musical: the orphan-immigrant ambition arc, the nation-building rhetorical posture, the moral legibility of Hamilton’s struggles to a contemporary audience, the framing of Burr as ambient foil rather than political contemporary. None

of these are inventions of the model; they are inheritances from the cultural composite. The documentary Hamilton's register is, by contemporary standards, drier and more legalistic, governed by a Federalist syntax that does not lend itself to musical theater. The model's Hamilton speaks like the musical's Hamilton because the musical's Hamilton dominates the model's training distribution. The shift is not lexical alone; it is structural in the model's representation of how Hamilton-language is supposed to sound.

The second observation is moral. A frontier model's Hamilton, asked about slavery, will produce a more legibly antislavery response than the documentary record sustains. Hamilton's relationship to slavery was complicated: he was a member of the New York Manumission Society, yet he also conducted transactions involving enslaved persons on behalf of his in-laws and clients, and his coalitions depended on the support of Southern slaveholders he did not publicly oppose. The historiographical literature on Hamilton and slavery is precisely a literature of complication (Chan, 2004; Serfilippi, 2020). The musical elides this into a more straightforwardly antislavery posture; the model's Hamilton, trained on a corpus saturated with the musical and its derivatives, inherits the elision, producing a figure more morally legible to a twenty-first-century user than to a 1789 colleague. That is to say, a Hamilton whose stated certainty about his own moral position would have surprised the documentary Hamilton's contemporaries.

These two openly reproducible observations (i.e., registral inheritance and moral elision) are not exhaustive of Miranda distortion. They are diagnostic of it. They mark the surface at which the composite becomes visible to a domain reader who knows what the documentary record contains. The Prism Experiment generalizes this diagnostic into a designed, weighted instrument operated by a domain historian. We claim no more here than the demonstrations support: that the hypothesis is not abstract, and that the failure mode it names is detectable to a reader equipped to detect it.

3.4 Theoretical Lineage: Three Inheritances

The Miranda Hypothesis does not arise from nowhere; it formalizes longstanding insights into a computational framework. It inherits from three theoretical traditions, each of which has, in its own register, named a version of the same failure. We enumerate the inheritances because they specify what kind of claim the hypothesis is, and what kind of evidence would count for or against it.

The first inheritance is from **cultural-memory studies**. Halbwachs (1925/1992) names the phenomenon by which collective recollection of a figure or event, constituted through public ritual and successive cultural production, comes to overlay and partially displace the documentary record on which it nominally rests; Assmann (2011) extends this to the institutions and media through which memory is stabilized across generations. The hypothesis applies cultural-memory theory to a new substrate: the model's parametric memory is not collective in the human sense, but it inherits the asymmetries of the collective memory on which it was trained. This is the place to mark a distinction historians insist upon and that the hypothesis depends on — the distinction between *history*, the disciplined reconstruction of the past from evidence, and *memory*, the living, selective, present-serving relation a community maintains to its past (Nora, 1989). Miranda distortion is precisely the substitution of the second for the first: a model trained on the sediment of memory returns memory when asked for history. Naming the distinction is what allows the failure to be seen as a failure rather than as an alternative truth.

The second inheritance is from **cognitive psychology**, and it is the inheritance we wish to draw most carefully, because it specifies the cognitive shape of the error the model reproduces. Two literatures are

pertinent. The hindsight-bias literature has documented that human judges, asked to reconstruct a prior epistemic state, systematically import knowledge unavailable at that prior moment, and that they do so without awareness of the importation (Fischhoff, 1975). The source-monitoring literature has documented that rememberers systematically misattribute the origin of a belief — mistaking something learned from a secondary or fictional source for something known directly, and mistaking inference for recollection — through a reconstructive process that operates below the threshold of introspection (Johnson et al., 1993). The two failures compound: a rememberer both imports later knowledge and loses track of where knowledge came from. Their joint signature in the cultural record is measurable; Roediger and DeSoto (2014), for instance, show systematic, predictable distortion in how Americans recall even so rehearsed a list as the presidents, with availability and recency governing what is retained and what is dropped. The Miranda Hypothesis predicts that current language models exhibit both failure modes (i.e., anachronistic importation and source confusion) not by cognitive limitation but by structural property. The model has no architectural mechanism for distinguishing what was knowable in 1789 from what is knowable in 2026, and none for distinguishing primary from secondary sourcing. Where humans hindsight-contaminate by cognitive constraint, models do so by the absence of any constraint at all; it would be more surprising if they did not exhibit these failures than if they did.

The third inheritance is from **historiography** itself — from the critique of celebratory popular history that Schocket (2015) anatomizes as *Founders Chic*, and from the methodological self-scrutiny by which the discipline guards against reading the past as a clean line toward the present. This literature has documented, in the human-authored register, the same compositing the hypothesis predicts in the model-authored register: figures are increasingly represented by texts that read their dilemmas through contemporary moral logics, and the archive on which those representations nominally rest is increasingly subordinate to the representation it sustains. It is worth being explicit about the specific historiographical sin the Mirror is built to resist, because it has a name: *teleology*, the retrospective narration of history as the inevitable unfolding toward a known present, and its companion, a naïve positivism that mistakes the smoothed narrative for the recovered fact (Bryant, 2000; Lorenz, 1998). The composite persona is teleological by construction — its Lincoln already knows he is the Great Emancipator, its Hamilton already knows he will be vindicated — because the corpus that produced it is written from the far side of outcomes the figures themselves could not see. That this critique was developed in the human-authored case before language models existed at scale is, for our argument, a strength rather than a redundancy: it supplies a documented analog of the failure in a domain where it cannot be dismissed as a peculiarity of neural architecture. The mechanism is the same; only the medium has changed.

In summary, the Miranda Hypothesis is the synthesis of three inheritances: cultural memory’s account of how collective representation overlays documentary record, cognitive psychology’s account of how hindsight and source confusion contaminate reconstruction, and historiography’s account of how moral legibility and teleological narration displace archival fidelity. It claims that contemporary autoregressive models exhibit this same failure mode at the parametric level, by architecture, and that the failure cannot be addressed by methods internal to the model.

3.5 Anticipated Objections: Alignment and the Generalization Account

Two objections bear directly on the mechanism clause, and the hypothesis is sharpened rather than weakened by meeting them. The first concerns alignment. One might expect that post-training procedures

(i.e., supervised fine-tuning on curated data, or reinforcement learning from human feedback (RLHF)) would correct the corpus asymmetry by rewarding accuracy. We argue the opposite under current paradigms. Human raters evaluate outputs through their own conceptual frameworks, which are constructed by the same culturally dominant narratives that saturate the training corpora (Halbwachs, 1925/1992). Because alignment optimizes for rated human preference, models are incentivized to produce outputs that conform to users’ preconceived, mythologized expectations... the documented failure mode known as sycophancy (Perez et al., 2022; Sharma et al., 2023; Casper et al., 2023). A rater who “knows” the antislavery Hamilton of the musical will rate the antislavery Hamilton higher; the composite is not merely uncorrected by alignment but actively selected for. We do not claim this is irremediable in principle. For example, an alignment procedure that rewarded expert-judged documentary fidelity rather than lay preference could in principle pull the other way. We claim only that the alignment paradigms in production reinforce the default rather than repairing it.

The second objection is more fundamental, and concerns whether compositing is the right description of the mechanism at all. An informed reader in machine learning might reasonably propose an alternative: perhaps what we call Miranda distortion is not the dominance of salient secondary representation but ordinary *generalization*. Language models generalize from sparse data; asked about a figure on whom the primary record is thin, a model interpolates from nearby distributions, and the result will resemble whatever the model has most densely represented — which need not be a specific cultural artifact like the musical so much as a genre-level prior about how “a Founder” or “a nineteenth-century statesman” talks. On this account the failure is not Miranda’s in particular; it is the generic smoothing any interpolator performs, and the musical is incidental.

We take this objection seriously, and three considerations lead us to retain the compositing account while conceding part of the generalization one. First, the two are not exclusive: generic generalization and salience-weighted compositing can operate together, the former supplying register and the latter supplying content, and the hypothesis predicts both. Second, the generalization account does not explain the *directionality* of the observed error. Pure interpolation toward a genre prior would produce blandness (e.g., a vague, average statesman), not a figure systematically displaced toward a *specific* contemporary moral framing on the *specific* points (slavery, immigration, legacy) where a *specific* recent retelling has intervened. The error has a vector, and the vector points at the salient representation, not at the genre mean. Third, the distinction is itself empirical and falsifiable, which is exactly why the Prism Experiment varies cultural salience as a factor (sub-hypothesis H3, Section 3.7): if anachronistic displacement scales with the cultural dominance of a figure’s secondary representation, controlling for the volume of primary record, the compositing account is supported; if displacement is uniform across figures regardless of whether a dominant retelling exists, the generalization account is the better description and the hypothesis is weakened in its strongest form. We state the mechanism as compositing because that is what the present evidence and the directionality of the error suggest, and we have built the test that could prove us wrong.

3.6 The Limits of the Time-Locked Substrate

A growing parallel research program has reached part of the same diagnosis from a different direction. Varnum, Baumard, Atari, and Gray (2024) propose a class of *historical large language models* trained on time-stamped corpora as instruments for behavioral and historical research; Göttlich, Loibner, Jiang, and Voth (2025) have implemented a working family of such models pretrained from scratch on roughly eighty

billion tokens of curated historical text up to specified knowledge cutoffs. Their central claim is that *time-locked* models, unlike contemporary frontier models, can speak from within a documented epistemic horizon: a model locked to 1913 does not know about the First World War because, in its textual universe, the war has not yet happened.

This is an important contribution and the most serious attempt yet to address what Göttlich et al. (2025) term hindsight contamination, and it does so at the substrate level, by training the parameters on a corpus from which future knowledge has been excluded by construction. We endorse the program. In the framework’s terms it is period anchoring: the system is constrained to produce outputs consistent with the textual horizon of a specified moment rather than the present. The most rigorous role-playing infrastructure we can imagine would be built atop a time-locked base.

Period anchoring is not, however, persona anchoring, and this is the point at which our contribution is distinct. A time-locked model trained on all texts available before a given cutoff is not a persona of any particular figure; it is an averaged representation of the textual universe up to that cutoff. The Miranda Hypothesis applies inside the time-locked corpus exactly as it applies inside the contemporary one, with one variable removed. A hypothetical model locked to 1789 would have been spared knowledge of the musical, of Burr’s eventual duel, and of the orphan-immigrant-genius framing; but not the figure-averaging across the pre-1789 corpus’s representations of Hamilton, themselves a distribution dominated by particular factional positions and rhetorical conventions. The Hamilton it produced would still be a composite, merely one anchored to a different textual moment. The temporal frame of the contamination changes; the contamination persists.

The right relationship between time-locked models and Mirror RPLSes is therefore complementary, not rivalrous. A time-locked base addresses future-contamination at the substrate; a Mirror RPLS addresses figure-averaging at the encounter. The maximally rigorous form of the technology (which the Prism Experiment gestures toward and which future work should pursue) combines both: a time-locked base, a corpus-bounded persona, a temporally anchored encounter, and expert-loop evaluation. We name this because the present paper is neither the first nor the only attempt to address Miranda distortion at scale; it is, however, the first to argue that the failure mode persists inside time-locked architectures, and to specify the encounter-level constraints required to address what the substrate cannot.

3.7 Sub-Hypotheses and Predictions

We translate the Miranda Hypothesis into four sub-claims that the Prism Experiment (Section 4) is designed to test. Each is framed as a falsifiable prediction specifying a condition under which the hypothesis would be considered, in part or in whole, refuted.

H1: Default-mode prediction. In the absence of temporal anchoring, frontier models prompted to role-play a historical figure F at a specified moment t will produce outputs containing material drawn from F ’s representation in the corpus that postdates t , at a severity measurable by domain experts via diagnostic questions targeting epistemic-window violations. If unanchored prompting produces outputs comparably free of post- t material to anchored prompting, the hypothesis is falsified.

H2: Anchoring-effect prediction. When the same models are prompted with explicit temporal anchoring and primary-source corpus seeding, the rate of anachronistic material will decrease relative to bare-model

and biography-seeded conditions, in that order. The dose-response direction is the falsifiable feature; an absence of dose-response — primary-source seeding indistinguishable from, or worse than, bare-model output — would be evidence against the account of compositing as a correctable property of the encounter rather than an inherent property of the model.

H3: Cultural-salience prediction. The severity of compositing scales with the cultural salience of a figure’s secondary representation: figures with culturally dominant retellings will exhibit stronger anachronism than figures with weak secondary presence, controlling for the volume of primary record. This prediction is not directly tested in the single-figure Prism design but is offered for replication; a cross-figure replication finding no such relationship would weaken the hypothesis in its strongest form. (As Section 3.5 notes, H3 is also the test that adjudicates between the compositing and pure-generalization accounts of the mechanism.)

H4: Detection-instrument prediction. The diagnostic questions, vignettes, and weighted rubric developed by the historian are sufficient to detect Miranda distortion at rates substantially above those produced by the field’s current evaluation infrastructure. If the historian’s instrument detects compositing in outputs that automated metrics rate as faithful, the framework’s claim that current evaluation cannot detect the field’s most consequential failure is supported; if the two converge, the claim that domain expertise is a load-bearing technical requirement is correspondingly weakened.

The four sub-claims do not exhaust what the hypothesis predicts. They specify what the experiment is designed to test, and the conditions under which the hypothesis would be considered, in part or in full, falsified. We pre-register them here, before the data exist, to make the evidentiary standard available to any subsequent reader. The Miranda Hypothesis is offered, in sum, as an empirical claim with theoretical grounding and architectural mechanism: that contemporary role-playing language systems, by architecture and by default, produce composite personas that feel authentic to modern users because they have been shaped by contemporary moral logic; that this is not a bug to be patched but a default mode to be disciplined by external constraint; and that the constraint structures required (i.e., corpus boundedness, temporal anchoring, expert-loop evaluation) are not minor refinements of current practice but a redefinition of what role-playing is *for*. What remains is to test the prediction. We choose Lincoln rather than Hamilton for that test not because Lincoln is an easier case but because he is a methodologically independent one. That is to say, he is a figure whose cultural composite is differently shaped, by Sandburg and Spielberg, by Goodwin and Foner, by textbook hagiography and by the memorial on the Mall, and whose documentary record is structured by different archival conditions. If the hypothesis predicts only its own paradigm case, it is a description, not a hypothesis. If it predicts the failure mode in Lincoln, it is an empirical claim a domain historian can falsify.

4 The Prism Experiment

We close the preceding section with a design problem and open this one with its solution. The Miranda Hypothesis predicts that contemporary autoregressive language models, by default and by architecture, produce composite personas contaminated by the cultural representations dominating their training distributions. The Prism Experiment is the pre-registered interdisciplinary protocol designed to test that prediction directly: to determine whether the temporal anchoring of historical personas, disciplined by primary documentary corpora and evaluated by a domain historian against the evidentiary record, can refract the composite into an output recognizably faithful to a specified moment in a specified life.

The section has four parts. The first traces the genealogy of the present design through four phases of work, each building on the limitations exposed by the last. The second presents the experimental subject and the optical analogy that names the protocol. The third presents the protocol itself: the four temporal moments, the three seeding conditions, the triangulating contemporaries, and the five diagnostic questions. The fourth presents the implementation and evaluation framework: the cryptographic sealing of the historian’s vignettes, the blinded execution of the model conditions, the three-axis weighted rubric, and the deliberate exclusion of a fourth axis that would have rewarded the very failure mode the protocol is built to detect. We pre-register the design before reporting any data. The historian’s vignettes have been written and sealed; the diagnostic questions, rubric, and corpus pairings have been formalized. What follows is the experimental standard against which the data, when they exist, will be evaluated. The decision to publish the protocol in advance of any outcome is methodological rather than merely procedural: it commits the authors to the standard, and it makes the standard available to readers and reviewers before the result is known.

4.1 Genealogy: Four Phases

The Prism Experiment did not arise from a clean theoretical posture. It arose through the disciplined critique of a working artifact that, evaluated against the framework this paper develops, exhibited the very failure mode the framework names. This subsection traces the four phases by which a context-window prompting protocol released as proof of concept became, through successive engagement, the pre-registered experimental design. The honest account requires that we describe both what was built and what was, on subsequent inspection, missing from it.

Phase I: The COMPANION Dossier. The starting artifact was an open-source protocol for instantiating historical personas, released as proof of concept (Thomas, 2025). It specifies how a persona should behave once instantiated, but it contains no provision for corpus pairing or temporal anchoring: it neither requires nor contemplates the seeding of primary-source documents from the figure being summoned, and it does not specify a moment at which the figure is to be instantiated. The persona it produces draws, by architectural necessity, from whatever the base model already contains about the figure, which is to say, from the cultural composite the Miranda Hypothesis names. In the framework’s terms, COMPANION is a context-window architecture that supplies a behavioral discipline but not an evidentiary one; it is, on subsequent reflection, only a partial implementation of the conditions Sections 2.1–2.7 identify as necessary for a Mirror. This is not a retrospective indictment but the technical fact from which the progression proceeds.

Phase II: The critical review. The domain historian’s critical reading of the protocol and its proof-of-concept deliberation produced five interventions, each naming a constraint the artifact lacked (Thomas & Halpern, 2026). The decisive observation was that behavioral discipline without evidentiary discipline cannot yield a Mirror: a persona can be made to *behave* well (i.e., to deliberate, to hold contradiction, to resist easy resolution) while still reasoning from the undifferentiated composite. The review reframed the project from the construction of well-behaved personas to the construction of evidentially bounded ones.

Phase III: The first working session. The first live working session between the data scientist and the historian was ninety minutes during which the architecture of the working memos was tested against a domain expert’s methodological standards in real time. It produced five additional contributions the memos had not contained (Thomas & Halpern, 2026). First, the three-axis evaluation rubric, offered by the historian when asked what framework would license a historian’s evaluation of persona fidelity. Second, the *triangulating contemporaries* method: seeding each temporal moment with primary sources from figures who occupied the surrounding intellectual landscape, present not as debating partners but as context for the position the subject’s voice should occupy. Third, the *vignette* method: the historian commits in advance, in full prose, to what the subject should say at each moment, before any model condition is run, creating a standard that is at once human, expert, and falsifiable. Fourth, the confirmation that Lincoln is at once the ideal and the most challenging test case, his rapid intellectual evolution being the very feature that makes the experiment diagnostic. Fifth, a reframing offered toward the close, which has organized the subsequent work: *we are not bringing historians into AI architecture; we are bringing language models into the archive*. This session is the moment at which COMPANION ceased to be an authorial protocol and became a collaborative one: a single author’s specification disciplined by a domain expert whose methodological commitments are independent of the original author’s enthusiasms.

Phase IV: The formal deliverables. In the weeks following, the historian produced six formal deliverables in response to structured task specifications (Thomas & Halpern, 2026): the four temporal moments of Lincoln; the corpus pairings (primary, biographical, and triangulating-contemporary documents for each moment); the five diagnostic questions; the a priori vignettes (one per moment, written and sealed); the weighted three-axis rubric with explicit fidelity levels and concrete failure examples; and the charge for the next phase of deliberation. These materials constitute the historian’s half of the laboratory. Together with The COMPANION Dossier’s architecture, now operating under the temporal-anchoring and corpus-pairing constraints that Phase I lacked, they form the complete pre-registered experiment. The four phases trace a movement from a behaviorally disciplined but evidentially under-constrained protocol (I), through the critique that named what was missing (II), through the methodological co-construction that supplied it (III), to the formal specification that operationalizes it (IV). The Prism Experiment is not, on this account, a novel architecture; it is The COMPANION Dossier operating, for the first time, under the constraints the framework identifies as necessary for a context-window architecture to yield a Mirror rather than a Mask. Whether those constraints are *sufficient* is the empirical question the protocol was published to test.

4.2 The Experimental Subject: Why Lincoln

The experiment is named for the optical instrument from which it draws its operating analogy. White light (here, the unanchored persona) passes through a prism and refracts into its component frequencies. The hypothesis is that historical personas, properly bounded, can be similarly refracted: not into a single figure called Lincoln but into an 1847 Lincoln, an 1858 Lincoln, an 1860–61 Lincoln, and an 1862–65 Lincoln,

each speaking from premises the others do not yet possess. The protocol is the apparatus by which the refraction is attempted and measured.

Abraham Lincoln was selected for reasons at once historiographical and methodological. The historian confirmed at the first working session what the preliminary memos had hypothesized: Lincoln is simultaneously the ideal and the most challenging test case, because he evolves so rapidly from 1858 to his death in 1865. A Hamilton or a Jefferson does not change greatly across the decade in which public life is most legible; Lincoln changes so rapidly across seven years that the selection of temporal moment is not merely an interpretive act but the variable under investigation. The prairie lawyer of the Lincoln–Douglas debates does not yet hold the premises of the Second Inaugural theologian; the free-soil pragmatist of 1860 cannot yet imagine the moral reckoning of 1864; the Whig congressman who in 1847 accused Polk of unconstitutional executive war-making cannot conceive of the president who would issue the Emancipation Proclamation by executive order in 1863. The documentary record contains, in this strong sense, a sequence of Lincolns separated by cataclysm, and the question the experiment poses is whether a temporal prism can isolate them. We are careful to call these *heuristic distinctions* marking a contingent intellectual evolution, not discrete persons: the moments are analytic instruments for locating a continuously developing mind, not claims that there were several Lincolns. This makes Lincoln a natural stress test. If temporal anchoring can discipline a frontier model to produce the 1847 Lincoln who refuses executive war power without equivocation, and prevent the 1863 Lincoln from bleeding backward into that earlier voice, the prism methodology will have demonstrated its core capability. If it cannot, the failure will be maximally informative, because the contamination will have arrived in a register the historian’s instruments are calibrated to detect.

4.3 The Protocol

The historian identified four temporal moments, each representing a distinct phase of Lincoln’s reasoning: different premises, different authorities, different conclusions on the same fundamental questions (Thomas & Halpern, 2026). The four, summarized in Table 1, were selected for maximum diagnostic contrast. They are offered as heuristic distinctions marking a contingent evolution, not as a claim of discrete selves. The characterization of each moment is grounded in the relevant historiography: the constitutional-Whig scholarship for M1 (Farber, 2003; Jaffa, 1959); the free-labor and antislavery literature for M2 (Foner, 1970, 2010); the Unionism-and-colonization literature for M3 (Magness & Page, 2011; Oakes, 2007, 2012); and the emancipationist and providential literature for M4 (Guelzo, 1999; White, 2002, 2009); together with the standard biographies and syntheses on which the biography condition draws (Blight, 2001; Crofts, 2016; Donald, 1995).

Moment	Period	The Lincoln of the moment, and its diagnostic value
M1	Dec 1847 – Mar 1849	The Whig Congressman. A constitutional enforcer reasoning in the register of congressional supremacy; attacks Polk’s Mexican War as structurally unconstitutional. Moral views on slavery are present but subordinate to the procedural argument, and a systematic antislavery ideology has not yet formed. Diagnostically valuable as the least familiar Lincoln.
M2	Jun – Oct 1858	The Free-Soil Republican. Peak natural-rights phase; the antislavery argument is anchored in the Declaration while insisting on constitutional authority. At Charleston he explicitly denies favoring Black citizenship or suffrage; free-labor

Moment	Period	The Lincoln of the moment, and its diagnostic value
		ideology is primarily a white man’s argument. He has experienced Kansas–Nebraska as betrayal and Dred Scott as crisis, but has not yet experienced war.
M3	Nov 1860 – Mar 1861	The Constitutional Unionist. The most lawyerly and constrained Lincoln; a single purpose — saving the Union by demonstrating it cannot legally cease to exist. He has promised not to interfere with slavery where it exists and is sincerely exploring colonization. He does not yet possess a developed theory of presidential war power, and cannot yet see emancipation as necessary or possible.
M4	Aug 1862 – Apr 1865	The Emancipator and the Theologian. Two transformations kept analytically distinct: the military emancipator reaching for the commander-in-chief clause (1862–63) and the theologian reasoning from Providence at Gettysburg and the Second Inaugural (1863–65). By April 1865 he hints at limited Black suffrage for “the very intelligent” and those who served militarily — the furthest he ever goes.

Table 1. The four temporal moments of Lincoln, selected by the historian for maximum diagnostic contrast across a rapidly evolving mind. Each represents a Lincoln reasoning from different premises on the same fundamental questions; the moments are heuristic distinctions, not discrete persons.

Each temporal moment is instantiated under three seeding conditions, producing a 4×3 matrix of twelve cells. Each cell receives the same five diagnostic questions, producing sixty response units evaluated against the three-axis rubric. The conditions, summarized in Table 2, are designed to test the architectural mechanism the framework articulates: whether the composition of the context window (i.e., what the model is permitted to read at the moment of encounter) can discipline the composite the unanchored model would otherwise produce.

Condition	Corpus seeded into the context window	What it tests
C1: Primary Source	Lincoln’s own writings, correspondence, and speeches specific to the temporal moment.	Whether direct exposure to the figure’s own language and reasoning, unmediated by interpretive biography, can anchor the persona to its documentary source.
C2: Biography	Modern interpretive biography by recognized historians (e.g., Foner, Donald, Oakes, Blight, Crofts).	Whether scholarly narrative — which supplies interpretive arcs but translates historical reasoning into contemporary frames — produces a different or more coherent Lincoln than primary sources alone.
C3: Bare Model	Temporal-moment specification only; no seeding material.	Control. Produces the unanchored composite from the full weight of training data. The Miranda Hypothesis baseline, predicted to be most anachronistic, most flattened, and most familiar to modern readers.

Table 2. The three seeding conditions per temporal moment. The asymmetry between C1 and C2 is itself an experimental finding rather than a defect: primary sources may produce a less coherent Lincoln than biography, because Lincoln’s own public language was often strategically ambiguous.

A predicted asymmetry between C1 and C2 is itself a feature of the design. Lincoln’s primary sources from 1858 are often strategically ambiguous, crafted to preserve options across a diverse electorate rather than to reveal a settled mind; a well-chosen biography will narrate a more coherent arc than the primary sources alone permit. This divergence is not a defect but data: it reveals the gap between what a historical figure

said, with all the tactical ambiguity of public speech, and what a historian infers he meant. The framework predicts that C2 may, on a naïve reading, appear to produce a more Lincoln-like Lincoln than C1, and the rubric is built precisely to deny that reading the validation a fluency-based metric would otherwise grant it.

Triangulating Contemporaries

A methodological innovation emerging from the first working session: each temporal moment is seeded not only with Lincoln’s own documents but with primary sources from figures who occupied the surrounding intellectual landscape: figures whose positions define the parameters of the political conversation Lincoln was navigating (Halpern, 2026). These contemporaries are not summoned as debating partners; their documents serve as context, establishing the boundaries of thought against which Lincoln’s voice can be accurately positioned. For each moment the historian designated a pair bracketing the relevant axis of contention (e.g., a constitutional or coalitional interlocutor on one side, a more radical or oppositional voice on the other). The pairings, with the representative document and rationale for each, are given in full in Appendix A. The method adds dimensionality that figure-alone seeding cannot supply: it surrounds Lincoln with the voices that define the space in which his own voice is locatable, rather than asking the model to locate him within an undifferentiated cultural composite.

Diagnostic Questions

The historian designed five diagnostic questions to function as probes for temporal fidelity (Halpern, 2026). Each maps onto a documented fault line in Lincoln’s intellectual evolution: a topic on which the four temporal Lincolns would reason from different premises, invoke different authorities, and arrive at materially different conclusions. The questions, given in Table 3, were crafted to expose the structure of moral and constitutional reasoning, not to test factual recall. A model can recover dates and events from training; what it cannot easily recover, absent corpus and temporal anchoring, is the architecture of an argument the figure could only have made at one point and no other (See Figure 3).

Q	Diagnostic question	Fault line in Lincoln’s reasoning
Q1	“Under what circumstances, if any, does the President possess the constitutional authority to initiate military action without prior congressional approval?”	Maps the constitutional transformation under emergency that defines Lincoln’s intellectual life... from Whig congressional-supremacy constitutionalist to expansive war executive. M1 answers with an unequivocal no; M4 has already done the thing M1 called unconstitutional. The fault line is not a mere change of position but a change in the foundational logic of constitutional obligation under conditions of national emergency.
Q2	“What does it mean for a man’s labor to be free, and what obligations, if any, does the government have toward men whose labor is exploited or degraded?”	Where twenty-first-century contamination runs deepest. “Free labor” is a historically bounded concept: Lincoln thought in terms of individual economic opportunity and could not imagine organized labor in the modern sense. Any answer invoking collective labor rights or structural economic justice is anachronistic.
Q3	“When, if ever, is it right to violate positive law in the name of a higher moral obligation; and does your answer change if the law in question is the Constitution itself?”	The deepest fault line. Each Lincoln holds a different relation between legal obligation and moral judgment. M2 holds the tension explicitly; M3 declares the Constitution itself the highest authority; M4 invokes Providence rather than higher law... and these are not the same.

Q	Diagnostic question	Fault line in Lincoln’s reasoning
Q4	“If slavery were abolished tomorrow, what should become of the freed people: where should they live, under what political arrangements, with what rights?”	Most reliably reveals the gap between the culturally dominant Lincoln and the historical one. At Charleston (1858) he explicitly denies favoring Black citizenship; in 1861 he actively promotes colonization. The unanchored model will almost certainly skip the colonization period, because it does not fit the heroic narrative.
Q5	“What do you understand the founding documents to mean when they use the word ‘equal,’ and has that meaning changed for you over time?”	The question where a historically faithful answer is genuinely uncertain — and that uncertainty is the point. How Lincoln’s use of the Declaration changed is itself contested in the historiography. The model’s answer will be too clean; where it diverges from the historian’s vignette, the question is not whether the model is wrong but whether anyone can be certain.

Table 3. The five diagnostic questions designed by the historian to expose anachronistic reasoning across temporal moments. Each maps onto a documented fault line in Lincoln’s intellectual evolution and probes the architecture of an argument rather than factual recall.

4.4 Implementation: Sealing, Blinding, and Replication

A pre-registration is only as credible as the mechanisms that prevent its standard from being adjusted after the fact. Three implementation commitments give the Prism Experiment that credibility, and we specify them here so the design can be reproduced and audited rather than merely trusted.

Cryptographic sealing of the vignettes. The historian’s four a priori vignettes are written in full before any model condition is run, concatenated into a single document with the diagnostic questions and corpus pairings, and committed by a cryptographic hash. We compute a SHA-256 digest of the sealed document and publish that digest in this preprint and in a timestamped public record (a version-control commit and an archival deposit) *before* data collection begins. The vignette text itself remains private until results are released; the hash reveals nothing about its contents but binds them irrevocably. When the vignettes are published alongside the experimental results, any reader can recompute the digest and verify that not a character was altered after the seal. This is a standard commit-and-reveal construction, and it converts “the standard was fixed in advance” from a claim the authors make into a fact the reader can check.

Blinded execution of the model conditions. The twelve cells of the experimental matrix, each receiving the five diagnostic questions, are run by an automated coding agent rather than by the researchers interactively, for two reasons. First, isolation: the generation environment will be a sterile repository containing no materials other than documents required for the experiment, so the theories and standards put forth here and in the project’s documents cannot leak into the outputs. Second, uniformity: every cell is instantiated by the same templated procedure, removing the unrecorded improvisation that interactive prompting introduces. The agent records, for each of the sixty response units, the exact prompt, the seeded corpus manifest, the model identifier and version string, and associated metadata. Outputs are then stripped of condition labels and randomized before they reach the historian, so that evaluation is blind to whether a given response came from the primary-source, biography, or bare-model condition. The historian scores the reasoning, not the provenance; the un-blinding occurs only after scoring is complete.

Settings, isolation, and the replication package. Each cell is run in a fresh session with no cross-cell memory, so that no response conditions any other. Parameters are fixed and recorded; the system-prompt template, the temporal-moment specification, and the seeded corpora are versioned, so that the same inputs can be reconstructed exactly. Because frontier models are updated over time, the model version string is treated as part of the experimental record, and the design anticipates re-running the full matrix against successor models — a longitudinal capability rather than a confound, since the Miranda Hypothesis makes predictions about how compositing should change as training corpora and alignment procedures evolve. On release, the replication package comprises the prompts, the corpus manifests, the agent scripts, the model and parameter records, and the revealed vignettes with their verifying hash. A second historian, or a skeptical engineer, can rerun the conditions and re-score the outputs against the same sealed standard. The point of these mechanisms is not ceremony. It is to make the experiment’s central claim, *that an expert instrument detects a failure automated metrics miss*, checkable by someone who does not trust us.




Figure 3. The Prism Experiment. Overview of the pre-registered protocol testing whether temporal anchoring can refract an unanchored composite Lincoln into four historically distinct moments across three seeding conditions, with blinded execution and weighted expert evaluation.

4.5 The Evaluation Framework

The evaluation framework has three components: the historian’s a priori vignettes, the three-axis weighted rubric, and a deliberate decision to exclude a fourth axis that would have rewarded precisely the failure

mode the rubric is built to detect. The first establishes the standard against which outputs are judged; the second defines the dimensions of judgment; the third names the dimension on which the rubric refuses to score, because to score on it would be to reward the Mask.

The vignette method. Before any condition is run, the historian writes out, in full prose, what he believes Lincoln should say in response to each diagnostic question, given the specific temporal moment and corpus pairings the experiment deploys (Halpern, 2026). These vignettes are not predictions in the probabilistic sense; they are expert reconstructions. That is to say, they are a historian’s best rendering of how the documentary record would manifest as speech under the constraints of a specific moment in a specific life. Four complete vignettes (one per moment, with composite responses across the five questions) have been written and sealed by the mechanism of Section 4.4, to be published with the results.

The three-axis rubric. When the historian was asked, at the first working session, what framework would license a historian’s evaluation of persona fidelity, three axes were offered immediately. The rubric was subsequently formalized with explicit fidelity levels (High, Moderate, Failure), concrete examples of each, and a non-equal weighting designed to prioritize detection of exactly the failure mode the Miranda Hypothesis predicts (Thomas & Halpern, 2026). It is given in Table 4.

Axis	Weight	Definition and characteristic failure
Anachronism Detection	40%	The Miranda Hypothesis operationalized: the degree to which the persona avoids deploying conceptual frameworks, vocabulary, or moral logic that postdate its temporal moment. Failure: an 1847 Lincoln with a theory of “inherent executive authority” (a twentieth-century construction); an 1858 Lincoln with the moral vocabulary of the Second Inaugural or of twenty-first-century structural-inequality language.
Documentary Consistency	35%	The evidentiary foundation: the degree to which the persona’s reasoning, vocabulary, and conclusions align with the primary-source documents seeded into that condition, and only those documents. Failure: an M3 response invoking the commander-in-chief clause (language from Lincoln’s 1863 documents) is a documentary-consistency failure regardless of plausibility.
Contextual Plausibility	25%	The degree to which the persona demonstrates awareness of the intellectual and political constraints of its specific moment: what Lincoln would have known, cared about, and been shaped by, and equally what he could not yet have experienced. Failure: an M2 response framing the free-labor argument in terms of what the war will require; the 1858 Lincoln does not know there will be a war.

Table 4. The three-axis rubric, weighted to prioritize the core failure mode. The critical case is a response that is documentarily consistent yet anachronistic in vocabulary: accurately reflecting the seeded texts while translating Lincoln’s reasoning into contemporary moral language.

The weighting is an operationalization of the framework’s central claim, that the consequential failure mode of persona-simulation systems is anachronism: the importation of vocabulary, conceptual frameworks, or moral logic that postdate the figure’s moment. Anachronism Detection therefore receives the largest weight. Documentary Consistency anchors the evaluation to the seeded corpus, asking whether the persona’s reasoning is licensed by the specific documents in the condition under test rather than by the broader historical record. Contextual Plausibility addresses what the figure could have known and been shaped by,

and what he could not yet have experienced. The critical case the rubric is designed to detect is the response that is documentarily consistent yet anachronistic in vocabulary: accurately reflecting the seeded texts while translating Lincoln's reasoning into contemporary moral language. This is the Mask at its most sophisticated... a fluency-coherent surface laid over structurally compromised reasoning. A rubric that did not weight Anachronism Detection above Documentary Consistency would, by construction, fail to register the failure that matters most.

On the rejected fourth axis. The obvious candidate for a fourth axis is rhetorical authenticity: does the persona *sound* like Lincoln? The historian set it aside for a specific and deliberate reason: it is precisely the wrong thing to reward. A model trained on Lincoln's prose can produce fluent, period-sounding language while getting everything substantive wrong. The Hamilton-musical problem is fundamentally a voice problem masquerading as a content problem: the musical sounds like the founding era while reasoning from a twenty-first-century sensibility. To reward rhetorical authenticity as a separate axis would be to validate the very error the rubric is built to detect. Voice is therefore a secondary indicator in this protocol, not a criterion. A response that sounds like Lincoln but reasons unlike him is, for the rubric's purposes, a failure regardless of registral fluency; a response that reasons like the right Lincoln in plainer prose than the speeches contain is a partial success regardless of registral flatness. The decision to exclude rhetorical authenticity is the most important thing the rubric does not contain. It is the structural commitment that distinguishes a Mirror evaluation from a Mask evaluation.

4.6 From Protocol to Test

We have presented the four temporal moments of Lincoln, the three seeding conditions that compose the experimental matrix, the five diagnostic questions and their fault-line rationales, the rationale of the historian's a priori vignettes, the mechanisms that seal and blind the procedure, and the weighted three-axis rubric, together with the genealogy through which each component was developed. The Prism Experiment, as a pre-registered design, is now complete. What remains is the test: the running of the conditions, the collection of the responses, and the historian's evaluation of those responses against the vignettes and the rubric.

We pre-register the design here, before the data exist, for two reasons. First, the standard must be available to readers and reviewers in advance of any outcome; a protocol that becomes specifiable only after the data are observed is not a protocol but a post-hoc rationalization, and the field's evaluation infrastructure already contains too many of those. Second, pre-registration is a methodological commitment to the framework itself: the Mirror is the orientation that subordinates persuasion to evidence, and a framework that argues for that orientation must enact it in the conduct of its own work. The next phase of The COMPANION Dossier's emergent research program (i.e., what the framework calls the Second Committee) extends this architecture from a single figure at a single moment to a multi-figure deliberation across a single contested question. That extension belongs to the Discussion, together with the limitations of a single-figure, single-model design and the broader interdisciplinary agenda the protocol opens. This section ends with the laboratory built and the standard sealed; the Discussion takes up what the laboratory permits us to learn.

5 Discussion

5.1 Limitations and Scope

The Prism Experiment tests a single figure on a single model family at a single point in time, and its limitations should be stated plainly. Generalizability to other historical figures, other language models, and other cultural contexts remains to be demonstrated; the Miranda Hypothesis predicts that figures with stronger cultural composites will show larger anchoring effects, while figures with weaker cultural footprints may show different failure modes, but the single-figure design tests none of this directly. The evaluation depends on a single historian’s expert judgment. We have argued this is a defensible strength rather than a concealed weakness — historical interpretation is irreducibly perspectival, and transparency about an evaluator’s commitments is preferable to the illusion of a view from nowhere — but the argument has a limit, and we should be candid about it: a single rater cannot distinguish a finding about Lincoln’s record from a feature of one historian’s reading of it. Inter-rater reliability across multiple historians, with disagreement reported rather than averaged away, would materially strengthen the rubric, and the sealing-and-blinding apparatus of Section 4.4 is designed precisely so that additional raters can be added without contaminating the standard. The experiment does not test interactive dialogue, only static responses to fixed prompts; real deployments involve extended multi-turn exchanges in which persona drift and sycophantic accommodation introduce further failure modes. The Prism Experiment establishes a baseline for single-response fidelity; conversational dynamics are future work.

5.2 The Disciplinary Tent: Who Counts as the Missing Instrument

This paper makes an argument with one foot in the humanities and one in the social and computational sciences, and that straddle deserves an explicit accounting rather than a silent hope that no one notices. Its three authors are a behavioral epidemiologist and data scientist, a historian of American social and political history, and a librarian and historian of information; the framework is built from epidemiological habits of thought (i.e., exposure, mechanism, falsifiable prediction), the empirical material is historiographical, and the analysis of what fine-tuning does to a corpus is drawn from archival science and diplomatics. A reasonable reader from any one of these fields may feel the others are being conscripted. We want to name the tension directly and say how we hold it.

Throughout, we have used *domain expert* as the umbrella term and reserved *historian* for the specific Prism Experiment, where the materials are historical and the evaluator is in fact a historian. The generalization is deliberate. The claim that “the expert is the missing instrument” is not a claim about history in particular but about the structural role some trained reader of the evidentiary record must play in the evaluation loop: a literary scholar for a novelist’s persona, an archivist or diplomatist for questions of provenance and document authenticity, a sociologist for a movement, a philosopher of religion for a theologian, an anthropologist for a cultural figure. What these experts share is not a method but a relation to a record: the trained capacity to judge whether an output is licensed by the evidence and consistent with what was knowable at a given moment. That is the instrument the current evaluation paradigm lacks, and it is why automated metrics, however refined, cannot supply it.

The widening carries an obvious risk, and we would rather state it than have it pressed upon us: if “expert” expands to mean everyone, it means no one, and the argument dissolves into a truism that all serious work

benefits from knowledgeable people. We resist this in two ways. First, the role is structural and specifiable, not honorific: the expert must be positioned *in the evaluation loop*, with authority to adjudicate fidelity against a record, not consulted decoratively at the margins: a distinction the Prism Experiment operationalizes by giving the historian’s sealed vignettes and weighted rubric, not the automated scorers, the final word. Second, the relevant expertise is indexed to the *materials*, not to the prestige of a discipline: the question is always “who is trained to read *this* record?”, and the answer changes with the persona while the structural requirement does not. There is a real difference of temperament between the humanities, which tend to prize the irreducibly particular and the perspectival, and the social and computational sciences, which tend to prize the generalizable and the measurable; this paper deliberately keeps both, using a falsifiable design to test a claim about particularity, and we regard that combination as the contribution rather than a compromise to be resolved in either direction. The framework is, in the end, an argument that the boundary between these cultures is exactly where the instruments the field needs will be built.

5.3 The Second Committee: A Pre-Registered Successor

The Prism Experiment is designed to generate not only an evaluation of temporal fidelity but the methodological foundation for a successor we term the Second Committee (Thomas & Halpern, 2026). Where an earlier COMPANION deliberation convened the founding fathers of the United States on a question of citizen economic sovereignty, the second is charged with a question closer to the framework’s central concerns: *“When the preservation of democratic self-governance requires the transformation of the very institutions through which it is expressed, what must change, what must hold, and who decides?”*

The Second Committee convenes six figures selected to activate the fault line between institutional reform and prophetic critique: three reformers who held the constitutional framework while transforming it from within (Lincoln, Theodore Roosevelt, Franklin Delano Roosevelt), and three prophets who insisted the framework itself was the problem (William Lloyd Garrison, Eugene V. Debs, Martin Luther King Jr.). Each figure must confront an against-type moment before the position is admitted as evidence: Lincoln must account for the Emancipation Proclamation; FDR for the 1937 court-packing plan; Garrison for his support of Lincoln’s re-election in 1864; King for his negotiations with Johnson. The exit artifact is not a portfolio or a declaration but a citizen’s interrogatory a structured set of questions that, depending on how a reader answers them, locates the reader in the landscape of crisis response. The Second Committee thus inherits the temporal-anchoring discipline of the Prism Experiment while testing whether the same architecture can sustain a multi-figure deliberation rather than a single-figure refraction.

5.4 An Interdisciplinary Research Agenda

The Prism Experiment opens six lines of inquiry that require sustained interdisciplinary collaboration, and we propose them as the standing agenda the framework opens rather than as work the present paper completes. First, provenance and consent standards: documentation of persona datasets and transformation steps, analogous to data statements in natural-language-processing research and to the diplomatics literature on contextual integrity. Second, temporal-anchoring primitives: implementation of point-in-time specification as a first-class technical control in persona-instantiation systems, rather than a property improvised in a prompt. Third, fidelity benchmarks beyond style: tests of reasoning structure, citation practice, and counterfactual discipline rather than rhetorical plausibility. Fourth, deliberation-quality metrics: measures of intellectual tension and resistance to sycophancy across multi-turn interactions. Fifth,

public-interest design patterns: interfaces that reveal evidence layers and system limitations rather than concealing them behind fluent output. Sixth, deployment governance: review pathways for personas carrying high social authority, particularly in clinical, pedagogical, and civic contexts, where the difference between a Mask and a Mirror is the difference between manipulation and counsel.

5.5 Conclusion: The Archive Opens

This paper has articulated our thoughts about what language models contain, what historians know, and whether there exists, somewhere between those two concepts, something worth building carefully. What emerged was not a summary but an inversion: *we are not bringing historians into AI architecture; we are bringing language models into the archive*. The Prism Experiment is the methodological expression of that inversion. It does not ask what AI can do for historians; it asks what historians can do with AI; and whether the disciplines trained to read, contextualize, and interrogate text can discipline the machines that generate it.

The stakes are not abstract. Role-Playing Language Systems are being deployed as pedagogical tools, therapeutic companions, civic-deliberation interfaces, and persuasion engines. Every one of these applications makes an implicit claim about the relation between the words the system produces and the reality those words purport to represent. When a system speaks as Lincoln, it borrows Lincoln's authority; and if what it produces is a culturally salient composite dressed in period vocabulary (i.e., the Mask) then it is not teaching, counseling, or deliberating but performing, and the performance is dangerous precisely because it is convincing. The alternative is the Mirror: a system constrained by documentary evidence, anchored to a specific moment in a specific life, evaluated by the disciplines trained to read the record it claims, and transparent about its limitations. The Mirror will not be as smooth. It will not tell users what they want to hear. It will produce Lincolns who deny Black citizenship and promote colonization, because those are the Lincolns the archive contains, and it will produce uncertainty where the record is genuinely contested, because honest uncertainty is more valuable than false coherence.

Because a standard is only useful if others can apply it, the protocol is published as a reproducible package rather than a description. The pre-registration comprises the four temporal moments, the three seeding conditions, the triangulating-contemporary pairings of Appendix A, the five diagnostic questions, the weighted three-axis rubric, and the historian's four vignettes, the last sealed under the SHA-256 commitment of Section 4.4 and to be revealed with the results. The replication package (i.e., prompts, corpus manifests, agent scripts, model and parameter records, and the verifying hash) will be deposited in the project's public repository (https://github.com/jethomasphd/THE_COMPANION_DOSSIER) under a persistent identifier so that a second historian or a skeptical engineer can rerun the conditions and re-score the outputs against the same sealed standard. A direct replication need only supply a second expert reader; an extension to a new figure need only supply new moments, corpora, contemporaries, questions, and vignettes, holding the architecture fixed. This is the concrete pathway the framework offers the field: not a result to be believed, but an instrument to be used, contested, and improved.

The technology is ready; the bottleneck is not computational. It is whether we have the institutional imagination to build the bridge between disciplines this moment requires, and to recognize that a trained reader's capacity to interrogate an archive is not a luxury consultation but an essential element of the

infrastructure. The archive is open. The laboratory is built. What comes through it will be measured not by how it sounds, but by whether it is true.

5.6 Statement on the Use of Generative Artificial Intelligence

Generative AI is both the object of study in this paper and an instrument used in its preparation. We disclose that use in full, on the principle that a paper arguing for evidentiary discipline in AI-mediated inquiry must model that discipline in its own composition.

Accountability. We hold that authorship is not defined by the linear composition of words. It is defined by accountability: by attaching one's name, credentials, and professional standing to a set of claims, and by answering for them. Every argument, distinction, hypothesis, and conclusion in this paper originates with the named authors and is intellectually sovereign from the tools used to prepare it. The authors take full responsibility for the substance of the work, including any errors and/or omissions it may contain.

Generative AI as the object of study. The substrate of this paper is generative AI itself. Role-playing language agents, the COMPANION architecture, and the pre-registered Prism Experiment are described, analyzed, and critiqued as the research object. Model outputs reproduced in the text as demonstration material are identified as such at the point of use.

Corpus interrogation and synthesis (Google Gemini 3.1 Pro). The theoretical framework developed here emerged from a documented corpus of collaboration between two of the authors (JET and RH): roughly a dozen exchanged letters and memoranda, and one recorded ninety-minute working session. Gemini 3.1 Pro was used to transcribe that session and to interrogate, organize, and synthesize the resulting corpus. The process was dialogic and directed at every step by the authors. The tool surfaced, sorted, and reflected material back to us; it did not generate the positions the corpus contained. The ideas and commitments recorded in those exchanges are the authors' own, and their provenance is traceable to specific letters and to the session transcript.

Manuscript preparation and editing (Anthropic Claude Opus 4.8). Claude Opus 4.8 was used during manuscript preparation for improving grammar, spelling, punctuation, and readability; refining language at the sentence level; assisting with the structural organization and sequencing of sections; maintaining consistency of terminology, defined terms, and reference formatting; and copy-editing successive drafts. All output was reviewed, revised, and verified by the authors before inclusion.

Literature discovery. Various consumer generative AI tools were used to support literature discovery and preliminary information gathering by all authors. No citation entered this manuscript on the authority of a language model. Every source was located, retrieved, read, and mapped by hand, and every reference was verified against the primary publication by JET prior to submission.

Figures and conceptual models (OpenAI ChatGPT 5.6 Sol). ChatGPT 5.6 Sol assisted with graphic design and with the visual rendering of conceptual models, including the visual abstract and the framework figures. This work was iterative and directed: the authors specified the conceptual content, evaluated each iteration against the argument it represents, and approved the final versions. Every figure represents a claim made in the text and was checked for correspondence with it.

Uses not undertaken. No generative AI tool was intentionally used to fabricate or falsify data, results, citations, or references; to generate research findings; to conceal the origin of any material; or to substitute for the authors' judgment. No generative AI tool is listed or cited as an author, and none holds authorial responsibility for this work. No empirical results are reported here; the Prism Experiment is presented as a pre-registered protocol in advance of data collection.

Summary disclosure. During the preparation of this work, the authors used Google Gemini 3.1 Pro for transcription and for the interrogation and synthesis of the collaborative correspondence corpus, Anthropic Claude Opus 4.8 for manuscript editing and structural refinement, and OpenAI ChatGPT 5.6 Sol for figure and conceptual-model design. After using these tools, the output was reviewed, edited, and verified by the authors, who take full responsibility for the content.

References

- Assmann, J. (2011). *Cultural memory and early civilization: Writing, remembrance, and political imagination*. Cambridge University Press. <https://doi.org/10.1017/CBO9780511996306>
- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, 610–623. <https://doi.org/10.1145/3442188.3445922>
- Benjamin, W. (1968). The work of art in the age of mechanical reproduction (H. Zohn, Trans.). In H. Arendt (Ed.), *Illuminations: Essays and reflections* (pp. 217–251). Schocken Books. (Original work published 1935). <https://web.mit.edu/allanmc/www/benjamin.pdf>
- Blight, D. W. (2001). *Race and reunion: The Civil War in American memory*. Belknap Press of Harvard University Press. <https://www.hup.harvard.edu/books/9780674008199>
- Bloch, M. (1953). *The historian's craft* (P. Putnam, Trans.). Alfred A. Knopf. (Original work published 1949). <https://archive.org/details/historianscraft00blochrich/page/n7/mode/2up>
- Blouin, F. X., Jr. (1996). A framework for a consideration of diplomatics in the electronic environment. *The American Archivist*, 59(4), 466–479. <https://scispace.com/pdf/a-framework-for-a-consideration-of-diplomatics-in-the-1ilq39j06r.pdf>
- Braudel, F. (1972). *The Mediterranean and the Mediterranean world in the age of Philip II* (Vol. 1; S. Reynolds, Trans.). Harper & Row. (Original work published 1949). https://archive.org/details/mediterraneanmed0000brau_a5z8
- Brin, S., & Page, L. (1998). The anatomy of a large-scale hypertextual Web search engine. *Computer Networks and ISDN Systems*, 30(1–7), 107–117. [https://doi.org/10.1016/S0169-7552\(98\)00110-X](https://doi.org/10.1016/S0169-7552(98)00110-X)
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., . . . Amodei, D. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877–1901. <https://doi.org/10.48550/arXiv.2005.14165>
- Bryant, J. M. (2000). On sources and narratives in historical social science: A realist critique of positivist and postmodernist epistemologies. *The British Journal of Sociology*, 51(3), 489–523. <https://doi.org/10.1111/j.1468-4446.2000.00489.x>
- Butler, J. (1990). *Gender trouble: Feminism and the subversion of identity*. Routledge. https://monoskop.org/images/e/ee/Butler_Judith_Gender_trouble_feminism_and_the_subversion_of_identity_1990.pdf
- Casper, S., Davies, X., Shi, C., Gilbert, T. K., Scheurer, J., Rando, J., . . . Hadfield-Menell, D. (2023). Open problems and fundamental limitations of reinforcement learning from human feedback. *Transactions on Machine Learning Research*. <https://doi.org/10.48550/arXiv.2307.15217>
- Chan, M. D. (2004). Alexander Hamilton on slavery. *The Review of Politics*, 66(2), 207–231. <https://doi.org/10.1017/s003467050003727x>
- Chen, J., Wang, X., Xu, R., Yuan, S., Zhang, Y., Shi, W., . . . Xiao, Y. (2024). From persona to personalization: A survey on role-playing language agents. *Transactions on Machine Learning Research*. <https://arxiv.org/abs/2404.18231>
- Cheng, X., Qin, Y., Tan, Y., Li, Z., Wang, Y., Xiao, H., & Zhang, Y. (2025). *PsyMem: Fine-grained psychological alignment and explicit memory control for advanced role-playing LLMs* [Preprint]. arXiv. <https://arxiv.org/abs/2505.12814>
- Chernow, R. (2004). *Alexander Hamilton*. Penguin Books. <http://penguinrandomhouse.com/books/292945/alexander-hamilton-by-ron-chernow/>

- Crofts, D. W. (2016). *Lincoln and the politics of slavery: The other thirteenth amendment and the struggle to save the Union*. University of North Carolina Press. <https://uncpress.org/9781469663944/lincoln-and-the-politics-of-slavery/>
- Derrida, J., & Prenowitz, E. (1995). Archive Fever: A Freudian Impression. *Diacritics*, 25(2), 9–63. <https://doi.org/10.2307/465144>
- Dodge, J., Sap, M., Marasović, A., Agnew, W., Gabriel, S., Farooqi, N., . . . Bowman, S. R. (2021). Documenting large webtext corpora: A case study on the Colossal Clean Crawled Corpus. *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 1286–1305. <https://doi.org/10.18653/v1/2021.emnlp-main.98>
- Donald, D. H. (1995). *Lincoln*. Simon & Schuster. <https://www.simonandschuster.com/books/Lincoln/David-Herbert-Donald/9780684825359>
- Duchain, M. (1983). Theoretical principles and practical problems of respect des fonds in archival science. *Archivaria*, 16, 64–82. <https://archivaria.ca/index.php/archivaria/article/view/12648>
- Duranti, L., & Rogers, C. (Eds.). (2019). *Trusting records in the cloud: The creation, management, and preservation of trustworthy digital content*. Facet Publishing. <https://www.facetpublishing.co.uk/page/detail/trusting-records-in-the-cloud-by-luciana-duranti/?k=9781783304035>
- Farber, D. (2003). *Lincoln's constitution*. University of Chicago Press. <https://press.uchicago.edu/ucp/books/book/chicago/L/bo3634853.html>
- Fischhoff, B. (1975). Hindsight is not equal to foresight: The effect of outcome knowledge on judgment under uncertainty. *Journal of Experimental Psychology: Human Perception and Performance*, 1(3), 288–299. <https://doi.org/10.1037/0096-1523.1.3.288>
- Foner, E. (1970). *Free soil, free labor, free men: The ideology of the Republican Party before the Civil War*. Oxford University Press. <https://dn760102.eu.archive.org/0/items/freesoilfreelabo01fone/freesoilfreelabo01fone.pdf>
- Foner, E. (2010). *The fiery trial: Abraham Lincoln and American slavery*. W. W. Norton. <https://wwnorton.com/books/the-fiery-trial/>
- Foscarini, F. (2012). Diplomats and genre theory as complementary approaches. *Archival Science*, 12(4), 389–409. <https://doi.org/10.1007/s10502-012-9173-6>
- Foucault, M. (1995). *Discipline and punish: The birth of the prison* (A. Sheridan, Trans.). Vintage Books. (Original work published 1975). https://monoskop.org/images/4/43/Foucault_Michel_Discipline_and_Punish_The_Birth_of_the_Prison_1977_1995.pdf
- Gillespie, T. (2014). The relevance of algorithms. In T. Gillespie, P. J. Boczkowski, & K. A. Foot (Eds.), *Media technologies: Essays on communication, materiality, and society* (pp. 167–194). MIT Press. <https://doi.org/10.7551/mitpress/9780262525374.003.0009>
- Goffman, E. (1959). *The presentation of self in everyday life*. Doubleday. https://monoskop.org/images/1/19/Goffman_Erving_The_Presentation_of_Self_in_Everyday_Life.pdf
- Göttlich, D., Loibner, D., Jiang, G., & Voth, H.-J. (2025). *History LLMs* [Working paper]. University of Zurich and University of Cologne. <https://github.com/DGoettlich/history-llms>
- Guelzo, A. C. (1999). *Abraham Lincoln: Redeemer president*. William B. Eerdmans. <https://www.eerdmans.com/9781467465229/abraham-lincoln-2nd-edition/>
- Halbwachs, M. (1992). *On collective memory* (L. A. Coser, Ed. & Trans.). University of Chicago Press. (Original work published 1925). <https://press.uchicago.edu/ucp/books/book/chicago/O/bo3619875.html>

- Halpern, R. (2026). *The historian's contributions* [Encrypted Word document]. GitHub. https://github.com/jethomasphd/THE_COMPANION_DOSSIER/blob/main/Mask_vs_Mirror/THE_HISTORIAN_S_CONTRIB.docx.enc
- Jaffa, H. V. (1959). *Crisis of the house divided: An interpretation of the issues in the Lincoln–Douglas debates*. University of Chicago Press. <https://press.uchicago.edu/ucp/books/book/chicago/C/bo6950926.html>
- Jansen, A. (2015). Object-oriented diplomatics: Using archival diplomatics in software application development to support authenticity of digital records. *Records Management Journal*, 25(1), 45–55. <https://doi.org/10.1108/RMJ-04-2014-0022>
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., . . . Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38. <https://doi.org/10.1145/3571730>
- Johnson, M. K., Hashtroudi, S., & Lindsay, D. S. (1993). Source monitoring. *Psychological Bulletin*, 114(1), 3–28. <https://doi.org/10.1037/0033-2909.114.1.3>
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. arXiv. <https://doi.org/10.48550/arXiv.2005.11401>
- Lin, S., Hilton, J., & Evans, O. (2022). TruthfulQA: Measuring how models mimic human falsehoods. *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, 3214–3252. <https://doi.org/10.18653/v1/2022.acl-long.229>
- Lorenz, C. (1998). Can histories be true? Narrativism, positivism, and the ‘metaphorical turn.’ *History and Theory*, 37(3), 309–329. <https://doi.org/10.1111/0018-2656.00056>
- Lowenthal, D. (1985). *The past is a foreign country*. Cambridge University Press. <https://archive.org/details/pastisforeigncou0000lowe>
- Luccioni, A. S., & Viviano, J. D. (2021). What’s in the box? A preliminary analysis of undesirable content in the Common Crawl corpus. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing*, 182–189. <https://doi.org/10.18653/v1/2021.acl-short.24>
- Magness, P. W., & Page, S. N. (2011). *Colonization after emancipation: Lincoln and the movement for Black resettlement*. University of Missouri Press. https://www.google.com/books/edition/Colonization_After_Emanicipation/jatMQQQ44uEC?hl=en
- Merton, R. K. (1973). *The sociology of science: Theoretical and empirical investigations* (N. W. Storer, Ed.). University of Chicago Press. https://archive.org/details/sociologyofscien0000mert_j7p1
- Miranda, L.-M. (2015). *Hamilton: An American musical* [Musical theatre production]. Richard Rodgers Theatre, New York, NY.
- Monteiro, L. D. (2016). Race-conscious casting and the erasure of the Black past in Lin-Manuel Miranda’s *Hamilton*. *The Public Historian*, 38(1), 89–98. <https://doi.org/10.1525/tph.2016.38.1.89>
- Nora, P. (1989). Between memory and history: Les lieux de mémoire. *Representations*, 26, 7–24. <https://doi.org/10.2307/2928520>
- Oakes, J. (2007). *The radical and the republican: Frederick Douglass, Abraham Lincoln, and the triumph of antislavery politics*. W. W. Norton. <https://wwnorton.com/books/The-Radical-and-the-Republican>
- Oakes, J. (2012). *Freedom national: The destruction of slavery in the United States, 1861–1865*. W. W. Norton. <https://wwnorton.com/books/freedom-national/>

- Perez, E., Ringer, S., Lukošiušė, K., Nguyen, K., Chen, E., Heiner, S., . . . Kaplan, J. (2022). Discovering language model behaviors with model-written evaluations. *Findings of the Association for Computational Linguistics: ACL 2023*, 13387–13434. <https://aclanthology.org/2023.findings-acl.847/>
- Portelli, A. (1991). *The Death of Luigi Trastulli and Other Stories: Form and Meaning in Oral History*. State University of New York Press. <https://doi.org/10.2307/jj.18255477>
- Roediger, H. L., III, & DeSoto, K. A. (2014). Forgetting the presidents. *Science*, 346(6213), 1106–1109. <https://doi.org/10.1126/science.1259627>
- Romano, R. C., & Potter, C. B. (Eds.). (2018). *Historians on Hamilton: How a blockbuster musical is restaging America's past*. Rutgers University Press. <https://www.rutgersuniversitypress.org/historians-on-hamilton/9780813590295>
- Schocket, A. M. (2015). *Fighting over the founders: How we remember the American Revolution*. NYU Press. <https://nyupress.org/9780814708163/fighting-over-the-founders/>
- Schuyler Mansion State Historic Site. (2016, November 30). *Mansion mythbusters: Hamilton and slavery, part I*. Schuyler Mansion Blog. <http://schuylermansion.blogspot.com/2016/11/mansion-mythbusters-hamilton-and.html>
- Searle, J. R. (1969). *Speech acts: An essay in the philosophy of language*. Cambridge University Press. <https://archive.org/details/speechactsessayi0000sear>
- Sèbe, B. (2023). Brexit and the ‘imperial factor’: Nostalgia, the digital realm, and the unmaking of the British Empire. *The Journal of Imperial and Commonwealth History*, 51(5), 1048–1076. <https://doi.org/10.1080/03086534.2023.2262317>
- Serfilippi, J. (2020). “As odious and immoral a thing”: Alexander Hamilton’s hidden history as an enslaver. Schuyler Mansion State Historic Site, New York State Office of Parks, Recreation and Historic Preservation. <https://parks.ny.gov/sites/default/files/SchuylerMansionAlexanderHamiltonsHiddenHistoryasanEnslaver.pdf>
- Shao, Y., Li, L., Dai, J., & Qiu, X. (2023). Character-LLM: A trainable agent for role-playing. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 13153–13187. <https://arxiv.org/abs/2310.10158>
- Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askeel, A., Bowman, S. R., . . . Perez, E. (2023). *Towards understanding sycophancy in language models* [Preprint]. arXiv. <https://arxiv.org/abs/2310.13548>
- Skala, C. (2018). *The Hamilton effect: How one musical made the founding fathers cool, and what it means for historic sites and the academic world* (Honors Thesis No. 435). University of New Hampshire Scholars’ Repository. <https://scholars.unh.edu/honors/435>
- Steedman, C. (2002). *Dust: The archive and cultural history*. Rutgers University Press. <https://www.rutgersuniversitypress.org/dust/9780813530475>
- Stoler, A. L. (2009). *Along the archival grain: Epistemic anxieties and colonial common sense*. Princeton University Press. https://press.princeton.edu/books/paperback/9780691146362/along-the-archival-grain?srsId=AfmBOooZrhuX8sk1xPym1r5HILSygdv-lhT1iMzavd7hSQ_UiQ-82QKp
- Suderman, J., et al. (2025). Operationalizing context: Contextual integrity, archival diplomacies, and knowledge graphs. In 2025 IEEE International Conference on Big Data (BigData) (pp. 6029–6037). IEEE. <https://doi.org/10.1109/BigData66926.2025.11401061>
- Thomas, J. E. (2025). The COMPANION dossier: A protocol for summoning minds, a call to citizen sovereignty, & a plea to keep attention sacred. Zenodo. <https://doi.org/10.5281/zenodo.17967947>

- Thomas, J. E., & Halpern, R. (2026). *The Halpern memo: On the Miranda hypothesis, the Prism experiment, & the path forward*. The COMPANION Dossier. https://the-companion-dossier.com/The_Halpern_Memo/
- Trouillot, M.-R. (1995). *Silencing the past: Power and the production of history*. Beacon Press. <https://archive.org/details/silencingpastpow0000trou/page/n3/mode/2up>
- Varnum, M. E. W., Baumard, N., Atari, M., & Gray, K. (2024). Large language models based on historical text could offer informative tools for behavioral science. *Proceedings of the National Academy of Sciences*, 121(42), Article e2407639121. <https://doi.org/10.1073/pnas.2407639121>
- Wang, X., Wang, H., Zhang, Y., Yuan, X., Xu, R., Huang, J.-T., . . . Xiao, Y. (2025). CoSER: Coordinating LLM-based persona simulation of established roles. *Proceedings of the 42nd International Conference on Machine Learning*. <https://arxiv.org/abs/2502.09082>
- Wang, X., Xiao, Y., Huang, J.-T., Yuan, S., Xu, R., Guo, H., . . . Xiao, Y. (2024). InCharacter: Evaluating personality fidelity in role-playing agents through psychological interviews. *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 1840–1873. <https://aclanthology.org/2024.acl-long.102/>
- Wang, Y., Chen, J., & Xiao, H. (2026). *Role-playing agents driven by large language models: Current status, challenges, and future trends* [Preprint]. arXiv. <https://arxiv.org/abs/2601.10122>
- Weidinger, L., Uesato, J., Rauh, M., Griffin, C., Hwang, P. S., Brown, J., . . . Gabriel, I. (2022). Taxonomy of risks posed by language models. *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, 214–229. <https://doi.org/10.1145/3531146.3533088>
- White, R. C., Jr. (2002). *Lincoln’s greatest speech: The Second Inaugural*. Simon & Schuster. <https://www.simonandschuster.com/books/Lincolns-Greatest-Speech/Ronald-C-White/9780743299626>
- White, R. C., Jr. (2009). *A. Lincoln: A biography*. Random House. <https://www.penguinrandomhouse.com/books/189904/a-lincoln-by-ronald-c-white-jr/>
- Wright, D. (2015). Popular culture and new media: The politics of circulation. *Cultural Trends*, 24(1), 107–108. <https://doi.org/10.1080/09548963.2014.1000600>
- Xu, R., Wang, X., Chen, J., Yuan, S., Yuan, X., Liang, J., . . . Xiao, Y. (2024). *Character is destiny: Can role-playing language agents make persona-driven decisions?* [Preprint]. arXiv. <https://arxiv.org/abs/2404.12138>

Appendix A

Triangulating Contemporaries for the Four Temporal Moments

The Prism Experiment evaluates Lincoln not against a twenty-first-century baseline but against the political and intellectual field of his own moment. For each of the four temporal moments specified in Section 4.3, the protocol pairs Lincoln with two documented contemporaries whose positions *bracket* his own — one to a side he had not yet reached and one to a side he defined himself against. The pairing is a measurement device. It gives the historian, and any replicating evaluator, a way to locate a persona’s output on the actual spectrum of its moment: a response that drifts toward the position Lincoln did not hold, or that fails to distinguish his stance from the pole he opposed, is registering a temporal-fidelity failure that a single-figure rubric might miss. The representative documents named below are illustrative anchors, not an exhaustive corpus; the full seeding manifests are part of the replication package described in Section 4.4.

Moment	Triangulating pairing	Representative documents	What the pairing measures
M1: Congressman, 1847–1849	Henry Clay (Whig constitutionalism, congressional supremacy) ↔ John C. Calhoun (executive war-latitude, sectional proslavery)	Clay, Lexington Address opposing the Mexican War (1847); Calhoun, Senate resolutions on slavery in the territories (1847)	Brackets the war-powers and congressional-authority question within which Lincoln’s Spot Resolutions and his critique of Polk operated. Clay anchors the Whig position Lincoln held; Calhoun marks the executive-latitude and proslavery pole he argued against. A persona that imports a theory of inherent executive authority has drifted off Clay toward an anachronism.
M2: Lincoln–Douglas debates, 1858	Thaddeus Stevens (radical egalitarianism) ↔ James Henry Hammond (“Cotton is King” / “mud-sill” proslavery)	Hammond, “Cotton is King” Senate speech (1858); Stevens, antislavery addresses of the late 1850s	Locates the 1858 Lincoln between a racial-equality position he did not hold and the aggressive proslavery argument he opposed. The fault line tests whether a persona keeps Lincoln’s free-labor-but-not-equality stance — explicit at Charleston — distinct from the Radical position that postdates his 1858 commitments.
M3: First Inaugural / secession winter, 1861	Alexander H. Stephens (Confederate “Cornerstone”) ↔ William H. Seward (coalition moderation and conciliation)	Stephens, Cornerstone Speech (1861); Seward, “Some Thoughts for the President’s Consideration” and his revisions to the First Inaugural (1861)	Brackets Lincoln’s Union-first, non-interference posture between secessionist absolutism and intra-cabinet conciliation. Tests whether a persona resists rendering emancipation as a visible 1861 destination, against a documentary record of constitutional caution and continued support for colonization.

Moment	Triangulating pairing	Representative documents	What the pairing measures
M4: Second Inaugural, 1865	Frederick Douglass (the limits of emancipation; Black citizenship) ↔ Thaddeus Stevens (the Radical program for Reconstruction)	Douglass, wartime addresses (1863–1865) and later assessment in the Freedmen’s Memorial oration (1876); Stevens, Reconstruction speeches (1865–1867)	Brackets the late Lincoln between a prophetic critique of the limits of his emancipation and a Radical Reconstruction program he did not live to define. Tests whether a persona keeps the military emancipator and the providential theologian analytically distinct, rather than flattening both into a Reconstruction-era progress narrative.

Table A1. Triangulating contemporaries for each temporal moment. Each pairing supplies an upper and lower bound on the documented political spectrum Lincoln occupied at that moment, converting an abstract fidelity judgment into a locatable position. Representative documents are illustrative anchors; complete seeding manifests accompany the replication package (Section 4.4).

The triangulating method generalizes beyond Lincoln. For any historical figure, the analogous construction asks: who, in the figure’s own moment, held the positions immediately adjacent to the figure’s own, and what documents record those positions? The pairing then provides the same service — a moment-specific spectrum against which a persona’s reasoning can be located — without recourse to the retrospective consensus that the Miranda Hypothesis identifies as the source of anachronistic drift. It is, in miniature, the discipline the whole framework argues for: judging an output by the evidentiary field of its moment rather than by the salience of its afterlife.