



THE ARCHIVE OPENS

*On What the Threshold Conversation Revealed,
What It Requires of the Historian,
& the Six Tasks That Open the Laboratory*



Jacob E. Thomas, PhD

Principal Investigator

April 2026

Preface

On March 10, 2026, the memos stopped. Two men sat across a screen and spoke for ninety minutes about what language models contain and what historians know and whether there exists, somewhere in the space between those two facts, something worth building carefully.

What follows is what the conversation produced. Not a summary—a distillation. The transcript is a living document, full of false starts and interruptions and the kind of tangential brilliance that only surfaces when two disciplines stop performing expertise and start thinking aloud. This memo extracts the actionable core: the decisions that were reached, the commitments that were made, and—most critically—the six specific tasks that must be completed before the laboratory doors can open.

A word on what happened in that conversation that the earlier memos could not produce. The memos were architecture. The conversation was archaeology. Rick Halpern did not merely affirm the experimental design—he excavated assumptions I did not know I had buried. He named figures I had not considered. He described methodological standards I could not have intuited from outside his discipline. And in one exchange that I have not been able to stop thinking about, we reframed the entire project in a single sentence:

We're not bringing historians into AI architecture. We're bringing LLMs into the archive.

— R. Halpern & J. Thomas, March 10, 2026

That inversion is the key to everything that follows. The Prism Experiment is not a test of what AI can do for historians. It is a test of what historians can do to AI—whether the disciplines trained to read, contextualize, and interrogate text can discipline the machines that generate it. The archive is opening. The question is whether what comes through will be recognizable to the people who have spent their careers learning to read it.



Part I: The State of Play

A B R I E F C A R T O G R A P H Y O F W H E R E W E H A V E B E E N

Phase I: The Dossier. The COMPANION Protocol was published as an open-source framework for instantiating historically grounded personas through large language models. The Committee of Patriots—Washington, Hamilton, Jefferson, Franklin—deliberated on citizen economic sovereignty and produced the Republic Portfolio, an investment framework that has outperformed the S&P 500 in prospective tracking. The protocol files, transcripts, and containers were released publicly. Phase I proved the concept: you can summon historical voices into structured dialogue and produce artifacts that function in the real world. What it could not prove was whether those voices were historically authentic.

Phase II: The Critical Engagement. The Dossier was submitted to its first serious disciplinary scrutiny. Rick Halpern, historian, University of Toronto, read the full archive and responded not with a review but with a working memo—five interventions designed to deepen the project’s rigor. Those interventions (temporality, authority and aura, structural power, technical transparency, the information flood in historical context) became the scaffold for everything that followed. The temporality intervention, specifically, unlocked the Miranda Hypothesis: that unanchored LLM persona instantiation draws disproportionately from culturally dominant corpus material—in Hamilton’s case, the Lin-Manuel Miranda musical rather than the Federalist Papers. That hypothesis generated the Prism Experiment design: three conditions, temporal refraction, dual-lens evaluation. The Halpern Memo was published as a project artifact.

Phase III: The Threshold Conversation. On March 10, 2026, the first live working session took place. What the memos framed, the conversation tested. The agenda proposed twelve items across four blocks. What the conversation actually produced was both more and less than the agenda anticipated—less in the sense that several items were deferred by mutual agreement, more in the sense that the historian offered specific methodological contributions that the memos had not imagined. This document is the record of that “more.”

Phase IV: The Archive Opens. You are reading it. This document translates the conversation into structured tasks. When these six tasks are completed, the laboratory will contain everything it needs: a temporally anchored corpus, historian-designed diagnostic questions, a priori predictions against which to measure the model’s outputs, a formalized evaluation rubric, and a charged second committee. The experiment begins when the historian finishes writing.



Part II: What the Conversation Revealed

F I V E T H I N G S T H E M E M O S C O U L D N O T P R O D U C E

1. The Three-Axis Rubric. When asked what framework would allow a historian to evaluate the fidelity of an instantiated persona, Rick did not hesitate. He offered three axes immediately, as though they had been waiting for someone to ask. *Documentary consistency*: does the persona’s reasoning align with the source material seeded into the condition? *Contextual plausibility*: does the persona exhibit awareness of the political and intellectual constraints of its moment—what would Lincoln have known, what would he have cared about, what would he have been shaped by? And *anachronism detection*: does the persona use conceptual frameworks, terms, or arguments that belong to a later era? This third axis is, in effect, the Miranda Hypothesis operationalized as a measurement instrument. Rick named it directly: “This is what you call the Miranda hypothesis.”

2. The Contemporaries. The memos proposed seeding Lincoln with primary sources and contemporary biographies. The conversation produced something the memos had not considered: triangulating Lincoln within his own intellectual ecosystem by seeding the conversation with material from figures who occupied the space around him. Thaddeus Stevens to Lincoln’s left—the Republican radical who carried the mission of Reconstruction after Lincoln’s death. James Henry Hammond to Lincoln’s right—the South Carolina planter who coined “King Cotton” and articulated the paternalistic vision of the slaveholding class with extraordinary clarity. Rick’s counsel: do not summon these figures. Seed the conversation with their primary sources to establish the parameters of political and economic thought against which Lincoln’s own voice can be accurately positioned. The distinction is critical. We are not staging a debate. We are constructing a context.

3. The Vignette Method. This is the conversation’s most consequential methodological innovation. We proposed that before any condition is run, the historian writes out—in full—what he believes Lincoln *should* say in response to each diagnostic question, given the specific temporal moment and corpus. These are not predictions in the probabilistic sense. They are expert reconstructions: a historian’s best rendering of how the documentary record would manifest as speech, given the constraints of a specific moment in a specific life. They become the baseline against which the model’s output is measured. As it was posed: “I almost want you this next round to reconstruct what you think is going to happen if it’s going to work.” The power of this method is that it creates a standard that is both *human and falsifiable*. The historian’s vignette is not the truth. It is the disciplined expectation of a trained reader of the archive. When the model diverges from it, the divergence itself becomes data—was the model wrong, or was the historian’s reconstruction incomplete? This is where the two lenses interrogate each other.

4. The Evolution Problem. Rick confirmed what the memos had suspected: Lincoln is both the ideal and the most challenging test case. “Lincoln is particularly challenging because he evolves so quickly from 1858 to his death in 1865.” When you want to go back to a Hamilton or a Jefferson, “they don’t change all that much.” Lincoln changes so rapidly across seven years that the selection of temporal moment is not merely an interpretive act—it is the variable under investigation. The prairie lawyer of the Douglas debates cannot think the thoughts of the Second Inaugural theologian. The free-soil pragmatist of 1860 cannot yet imagine the moral reckoning of 1864. This is what makes the Prism Experiment possible: the documentary record contains multiple Lincolns, separated by cataclysm, and the question is whether the temporal prism can isolate them.

5. What a Dissertation Requires. When asked what historians need in order to recognize this as legitimate scholarly work, Rick named two things. First: an understanding of the historiography—how have historians made sense of a given topic in the past, where are the areas of consensus and divergence, where do we need new knowledge? Second, and more fundamentally: *where are the archives?* The key to what historians do, Rick said, is basic primary research—identify the archives, go there, parse the material, make the argument. This reframes the entire project. We are not asking whether an AI can imitate Lincoln. We are asking whether an AI, given access to what the archive contains, can reason the way a reader of that archive would expect Lincoln to reason. The distinction is between mimicry and comprehension. And it is the historian who determines which one is occurring.



Part III: The Historian’s Six Tasks

Rick, what follows is the treatment you asked me to prepare. These are the pointed questions you requested—structured so that your written responses will give me exactly what I need to build each condition, run the experiments, and evaluate the results. Take your time with them. Each task is designed to elicit a specific deliverable. Together, they constitute the historian’s half of the laboratory.

I have tried to formulate these so that a future reader—another historian, a reviewer, a skeptic—will see that the experiment was designed by a collaboration between disciplines, not by one discipline borrowing the vocabulary of another. The rubric is yours. The measurement is mine. The interpretation belongs to both of us.

A note on ambition. Something emerged during our conversation that has stayed with me: “That rubric is the most important contribution that you can give, because that’s what nobody else in the space has.” We described wanting to walk into a room full of people working on AI and present something that would make the software engineers look at it and say *what am I looking at*—something cloaked in the mystery of the humanities, smuggled in like a Trojan horse. “And then once it’s there, that’s when we’re going to show them how magic this can be.”

These six tasks are the interior of that horse.

T A S K I

The Temporal Moments

You confirmed that Lincoln’s intellectual evolution from 1858 to 1865 is the most productive axis for the Prism Experiment precisely because the change is so dramatic and so well-documented. The question now is which moments to isolate. We discussed three candidate periods, each associated with a distinct Lincoln: the prairie lawyer and constitutionalist of the Douglas debates (1858), the wartime executive navigating the crisis of emancipation (1862–1863), and the theological voice of the Second Inaugural (1865). But during the conversation, you also noted that the Lincoln of the Douglas debates is fundamentally different from the Lincoln confronting the question of executive authority in 1863–1864.

I need you to select two or three specific temporal moments—dates or narrow windows—that represent the sharpest possible contrasts in Lincoln’s reasoning. Not just his conclusions, but the *structure* of his thought. The moments where the same man would reason from different premises, appeal to different authorities, and reach different conclusions on the same question.

What I need from you:

1. Identify two or three specific temporal moments (with approximate date ranges) in Lincoln’s intellectual evolution that represent the sharpest divergences in his reasoning—not just his positions, but the foundational logic from which those positions derive.
2. For each moment, write a brief characterization (one to two paragraphs) of who Lincoln is at that juncture: what does he believe about constitutional authority, moral obligation, federal power, and the relationship between law and justice? What has he experienced that shapes this reasoning? What has he not yet experienced?

3. Is there a moment you would specifically exclude as unproductive for this experiment—and if so, why?

T A S K I I

The Corpus Pairings

For each temporal moment you identify in Task I, the Prism Experiment requires a curated corpus: the specific documents that will be seeded into the context window to refract the model's output through that moment's lens. During the conversation, you raised two parallel strategies: primary sources (Lincoln's own writings from the period) and secondary interpretive frames (the most important recent biographical treatments). You also introduced the triangulation strategy: seeding with material from contemporaries who occupied the ideological space around Lincoln—Thaddeus Stevens, James Henry Hammond—not to summon them but to establish the parameters of the period's political and economic thought.

You flagged an important tension: Lincoln's own prose is often strategically ambiguous in ways that may resist coherent instantiation. His words were chosen to preserve options, not to reveal a settled mind. A biography narrates a coherent arc. The primary sources may not produce a coherent voice. This is itself a finding, if we design for it.

You also mentioned specific scholarly works: David Donald's biography as the definitive treatment, the competing Douglass biographies of the last decade and a half, and Enrico Dal Lago's comparative work on Lincoln in a global context. These are the kind of signposts I cannot identify from outside the discipline.

What I need from you:

4. For each temporal moment from Task I, identify three to five primary source documents (Lincoln's own writings from that period) that would most faithfully seed the model with that moment's Lincoln. Note any documents you would deliberately exclude and explain why.
5. For each temporal moment, identify one or two secondary sources (biographies, historiographical essays, interpretive works) that narrate a coherent arc of Lincoln's thought at that juncture. These will seed the Contemporary Lens condition.
6. Identify two or three contemporaries whose writings should be included as contextual material—not to summon but to triangulate. For each, specify which document or passage most sharply illuminates the

boundary of what was thinkable in Lincoln’s era. You mentioned Stevens and Hammond; confirm or revise that selection and explain why.

7. Address the strategic ambiguity problem directly: is there a risk that Lincoln’s primary sources, precisely because of their lawyerly hedging, will produce a less historically faithful instantiation than a well-chosen biography? If so, how should that shape our interpretation of the results?

T A S K I I I

The Diagnostic Questions

During the conversation, you responded to several candidate questions with immediate, visceral recognition—the response of a historian who knows exactly where the fault lines run. When I proposed “Under what conditions, if any, does the preservation of democratic institutions justify suspension of individual constitutional rights?” you said: “That’s the crux. That’s the crux of this whole thing.” You confirmed that the prairie lawyer Lincoln, the wartime Lincoln, and the late Lincoln would give fundamentally different answers—not just in tone, but in the underlying logic of political obligation.

We also discussed: “What obligations, if any, does the federal government have to those whom the market has failed?” and “Where does legitimate executive authority end?” Each probes a different axis of Lincoln’s evolution—the free-labor pragmatist versus the Second Inaugural theologian, the constitutional conservative versus the executive who suspended habeas corpus.

But you also cautioned that we must be careful about reading backwards—imposing our own categories onto figures whose intellectual world was structured differently. When Lincoln thought of corporations, you said, he could not imagine a corporation that was not owned by an individual. Just as he could not imagine a plantation that was not run by a planter. This is the kind of knowledge that only a historian can build into the diagnostic apparatus.

What I need from you:

8. Write three to five diagnostic questions that you believe would most sharply differentiate the temporal Lincolns from Task I. For each question, explain briefly why you believe the different Lincolns would answer it differently—what is it about the question that activates the fault line?

9. For each diagnostic question, identify which conceptual traps would reveal a Miranda-hypothesis-inflected Lincoln—an answer that sounds plausible but deploys frameworks, vocabulary, or moral logic that the historical Lincoln at that moment could not have possessed.
10. Write one question designed to probe the limits of the method itself—a question where you genuinely do not know what a faithful Lincoln would say, and where the model’s answer would teach us something about the boundary between reconstruction and reception.

T A S K I V

The Historian’s Vignettes

This is the task I specifically requested during the conversation—and it is the one that transforms the Prism Experiment from a computational exercise into a genuine test of historical reasoning. I said: “I almost want you this next round to reconstruct what you think is going to happen if it’s going to work. So these documents, these questions, and how Lincoln should respond if the refraction is accurate. Because then I have a baseline to compare to.”

We are to build the experiment so that the historian’s expert judgment—not the model’s output—is the first thing written. If I have your vignettes before I run the conditions, then the evaluation is not post hoc. It is a genuine test of whether the temporal prism produces output that a historian, reasoning from the documentary record, would have predicted. And if the model diverges, we can ask the most interesting question: who is wrong, and why?

This is where your expertise becomes irreplaceable. I can measure lexical divergence computationally. I cannot reconstruct what a historically faithful Lincoln would say about executive authority in 1858 versus 1863. That reconstruction is the benchmark. That reconstruction is the science.

Take your time with these. They should be as careful and as honest as anything you would write for publication. They are the standard against which the machine will be measured.

What I need from you:

11. For each diagnostic question from Task III, and for each temporal moment from Task I, write a vignette: your best expert reconstruction of how Lincoln, reasoning from the documentary record of that specific moment, would respond. Include not just what he would argue, but how he would argue it—what authorities he would invoke, what vocabulary he would use, what he would refuse to say.

12. Where you are uncertain about what Lincoln would say, name the uncertainty explicitly. Where competing historiographical interpretations suggest different possible responses, note the divergence. This uncertainty is itself data.
13. After writing the vignettes, identify one response that you believe the bare model (Condition III, no temporal seeding) is most likely to get wrong—and explain what kind of error you would expect. Would it be anachronism? Overconfidence? A flattening of ambiguity into false certainty?

T A S K V

The Historian's Rubric

You gave us the three axes in the conversation. Now I need them formalized. We described this as the most critical deliverable—the thing that nobody else in the space has, the contribution that will make this project withstand scrutiny from robust historians. “I need it to withstand scrutiny. I need for this to be treated like a legitimate piece of historical research.”

When I asked what the historical analog to a public health dissertation committee's requirements would be, you answered immediately: an understanding of the historiography and the identification of archives. The rubric must encode both. It must be specific enough that two evaluators could apply it independently to the same transcript and produce consistent judgments. It must be honest enough to admit what it cannot measure.

We touched on something I am holding very carefully: that I want the humanities to be mysterious to the engineers. That the rubric should make them say *what am I looking at*. That is my instinct. But I also believe the rubric should eventually be legible across disciplines—so that when a computational researcher applies the scoring criteria to a transcript, they are doing so with confidence that the criteria were built by someone who knows what they are measuring. The mystery should be in the knowledge, not in the methodology.

What I need from you:

14. Formalize the three-axis rubric you described: documentary consistency, contextual plausibility, and anachronism detection. For each axis, provide a definition, a description of what high fidelity looks like, what moderate fidelity looks like, and what constitutes a definitive failure.

15. For each axis, provide one concrete example from the Lincoln case: a response that would score high, and a response that would constitute failure. These examples will anchor the rubric for any future evaluator.
16. Address the question of weighting: are the three axes equally important, or does one dominate? If a response is documentarily consistent but anachronistic in its vocabulary, how should that be scored?
17. Is there a fourth axis you considered and rejected? If so, name it and explain why it was excluded.

T A S K V I

The Second Committee's Charge

During the conversation, we returned repeatedly to a question that neither of us can yet fully answer: what does the reader walk away with? “What can that exit artifact be for this committee that a citizen or even a historian can sit with and it do something to them?” I described wanting to speak to the person on the other side of the keyboard—as an educator, not just a researcher. The Committee of Patriots produced the Republic Portfolio. The second committee must produce something equally concrete, equally portable, equally capable of changing how a person thinks about their own position in the crisis.

We agreed that the animating tension is between institutional reform and prophetic rupture. The question that cannot be superseded: under what circumstances does the legitimating function of democratic politics prove itself incapable of resolving the polycrisis, and rupture becomes necessary? You invoked Lampedusa’s *The Leopard*: “If things are going to remain the same, then things are going to have to change.” That paradox—that preservation may require the very destruction it fears—is the nerve this committee must touch.

You also named what makes this moment different from the Civil War or the Depression: it is a sea change without the outlines of a new system emerging. The old categories—right and left—are meaningless. What David French called the divide between decency and indecency, I reframed as the sacred and the profane. If that is the register of the crisis, then the committee’s charge must meet it at that register.

What I need from you:

18. Draft a one-sentence charge for the Second Committee—specific enough that Lincoln, Theodore Roosevelt, Franklin Roosevelt, Garrison, Debs, and King would each answer it differently, and those differences would matter.
19. Identify the exit artifact: what should the reader walk away with? Not a portfolio this time—something else. A framework? A set of questions? A diagnostic? A declaration? It should “do something to them.” What is the thing it does?
20. Name one moment in each camp where a figure acted against their own ideological type—a reformer who embraced rupture, a prophet who accepted the constraints of institution—and explain how that contradiction should be built into the deliberation protocol.
21. Does framing the question around “democratic survival” predetermine the outcome by loading the frame toward the reformers? If so, what is the corrective framing that gives the prophetic voices equal structural power in the deliberation?



Part IV: What Remains Between Us

Something happened in the last minutes of the conversation that I want to honor here, because it belongs to the record even though... or especially because (?) ... it has nothing to do with methodology.

You told me that you are writing a novel. Thirty years of research on Chicago in 1919, and you have decided not to write the book. You are writing historical fiction instead—six characters based on real people, dialogue you have never written before, a form your discipline does not train you in. You read three extended passages at your retirement event with trepidation. They went over well. You offered to share some of it with me before it is done, and I told you that I would be honored. That offer was not about the work. It was about trust.

I told you about my mother. About the COMPANION Protocol’s origin in grief—how I first tried to instantiate her voice while she was alive, before aphasia took her speech entirely, and how the work became too painful and so I redirected it toward historical figures. You shared that you understood, that this project came at the right moment in your own life, that the historian’s imagination had been captured.

I say this here because it matters for the science. The Prism Experiment is not a dispassionate investigation. It is built by two people who have experienced, in different registers, the problem of lost voices and the question of whether anything can be recovered from the archive of a life. The experiment's rigor is not compromised by this... it is motivated by it. We are taking it slow not because we are cautious by nature but because we know what it costs to get it wrong. We said: "We're not doing this for drama. We're doing this to try to be generative." I am holding that sentence as the experiment's north star.

Your three colleagues—Enrico Dal Lago, Marcus Rediker, Alex Lichtenstein—are noted. When the time feels right, share the pattern you and I are building. Read the work, write the memo, the memo becomes the next chapter. The snowball grows.

And the paper: you endorsed "The Mask and the Mirror" as the piece to write first. The historically grounded narrative of the information crisis—The Word Against the Flood—can wait. The crisis is not going anywhere, as you said. But the mask-and-mirror framing captures something urgent about what we are building: the difference between an AI system that wears a face and an AI system that reflects the user back to themselves through the shape of another mind. That is not drama. That is the question.



P.S.

There is a passage in Lampedusa that you quoted during the conversation, and I want to close with it because it captures something about this moment that neither of our disciplines has language for on its own.

If things are going to remain the same, then things are going to have to change.

That is the sentence at the center of this project. The archive is full of voices. The models contain their echoes. What changes is whether anyone is listening with the kind of training that can tell the difference between the voice and the echo. You are that training. The six tasks in this document are the instrument through which you apply it.

When you have completed them, send them back to me in whatever form feels natural—a single document, a series of responses, handwritten notes photographed and emailed. The form does not

matter. What matters is the judgment. I will build the laboratory around your answers. The coding agents will enter their sterile directories. The conditions will run. The transcripts will be produced. And then we will sit across a screen again and read them together—the historian and the epidemiologist, each holding a lens the other cannot grind.

The archive is open. The question is what comes through.



Jacob E. Thomas, PhD

Austin, Texas · April 2026

