



THE HALPERN MEMO

*On the Miranda Hypothesis,
the Prism Experiment, & the Path Forward*



Jacob E. Thomas, PhD

Principal Investigator

February 2026

Preface

In February 2026, the COMPANION Dossier received its first serious critical engagement from a historian. The memo that arrived in response was generous, rigorous, and informed by a career spent studying how power is structured and contested in American life. It hardened the epistemic foundations of the entire COMPANION concept. This document is a record of what emerged in the space between disciplines: a historian and a behavioral epidemiologist, thinking together about what these AI machines can and cannot summon. The exchange illuminates the territory that opens when two fields collide over a shared problem and neither can solve it alone.



Part I: The Exchange

The initial outreach posed a question the Principal Investigator could not answer alone:

“Could this be a meaningful pedagogical or civic deliberation tool, or is it JUST sophisticated AI theater?” — J. Thomas, February 17, 2026

The letter described how the Committee of Patriots had deliberated on citizen economic sovereignty, how the resulting Republic Portfolio had been backtested against a decade of market data, and how it was outperforming in prospective tracking. It named the gap: the dialogue *sounded* like an argument frozen for two centuries and suddenly resumed. But sounding compelling is not the same as sounding true. Only a historian could evaluate the difference.

The response, when it came, was encouraging and precise:

“This is an important project and a very promising use of your skills. The ‘Committee of Patriots’ strikes me as a particularly rich way to stage encounters between historically grounded personas and contemporary problems, and I can absolutely see its potential as both a pedagogical and civic deliberation tool rather than ‘just’ AI theater.” — R. Halpern

Then the full Dossier was read. And a memo arrived—not a review but a working document, actionable interventions designed to deepen the project’s rigor. What follows synthesizes those interventions, the hypothesis they unlocked, and the experimental path forward.



Part II: The Critical Interventions

The memo opens by foregrounding a tension embedded in the choice of interlocutors: the small-r republican vision of the Founding Fathers was, in significant respects, profoundly anti-democratic. The Senate, the Electoral College, the indirect mechanisms of representation were structural expressions of suspicion toward popular rule. To summon these figures into our present is to invite assumptions about governance that the convener must make explicit.

From this frame, five specific interventions emerge.

Temporality

Historical actors were not static worldviews but evolving figures shaped by specific conjunctures. The protocol as designed summons a *name*. But Lincoln in 1858 is not Lincoln in 1863. Jefferson in private correspondence is not Jefferson in presidential office. The memo recommends temporal parameters and corpus-based anchoring—enabling the convener to specify not just *whom* to summon, but *when* in that person’s intellectual evolution.

Authority and Aura

The ritual language creates presence, and presence generates authority. When Hamilton “arrives” and deliberates, his speech carries a persuasive weight that derives not only from logical force but from the aura of Hamilton-as-Founder. The memo recommends embedding historiographical plurality directly into the protocol: pair figures with their later critics, design structured counter-voices, make the mediation visible rather than invisible.

Structural Power

The Republic Portfolio operates within markets. The memo asks the question the Founders cannot ask about themselves: is the proposed intervention fundamentally altering incentive architectures, or does it operate within the very logics it critiques? This is a blind spot built into the roster, not into the protocol.

Technical Transparency

Broader adoption among scholars and civic technologists requires a parallel exposition: model dependencies, reproducibility conditions, failure modes. The grimoire summons. A supplementary white paper must explain the machinery of summoning.

The Information Flood in Historical Context

Information crises are not new. The penny press, radio propaganda, television—each transition produced its own prophets and counter-technologies. What is genuinely unprecedented is the mechanism: algorithmic optimization and behavioral targeting at planetary scale. Situating the Flood within this longer arc strengthens the thesis rather than diluting it.



Part III: The Miranda Hypothesis

The temporality intervention unlocked something the Principal Investigator had sensed but could not name.

The large language models that power COMPANION are trained on massive text corpora. Within those corpora, the volume and recency of text associated with “Alexander Hamilton” is not dominated by Hamilton’s own writings. It is dominated by *Hamilton*: the Lin-Manuel Miranda musical. The cast recording, the libretto, the thousands of reviews, analyses, and cultural commentary generated after 2015. In raw corpus terms, Miranda’s Hamilton almost certainly outweighs the historical Hamilton by both volume and recency—the two factors that most influence how language models weight their outputs.

This is the **Miranda Hypothesis**: that the Committee of Patriots, instantiated through an unanchored LLM, may have produced not the Founders’ deliberation but *Miranda’s version of the Founders’ deliberation*.

Several observations support it. The Committee was notably sympathetic to the convener’s progressive orientation—more egalitarian, more suspicious of oligarchy, more attuned to the language of economic justice than the documentary record would strictly warrant. This tracks with Miranda’s interpretive lens. Hamilton’s voice in the session was performatively confident in a way that reads more as musical theater than as the cautious, legalistic prose of the historical Hamilton. And the ease with which all four personas engaged contemporary economic concepts may reflect the musical’s anachronistic fluency rather than a genuine fusion of 18th-century reasoning with modern context.

Paradoxically, the Miranda Hypothesis *validates* the protocol’s core claim. COMPANION genuinely instantiates a version of the named figure from the available corpus. The question is: **which corpus dominates?** Without temporal anchoring, the summoned persona is an unweighted composite shaped more by cultural salience than by documentary fidelity. The task now is to control which corpus shapes the vessel.



Part IV: The Prism Experiment

If unanchored instantiation draws disproportionately from culturally dominant sources, the solution is a *temporal prism*: a curated corpus seeded into the context window alongside the protocol files, refracting the model’s broader training data through a specific historical lens. The analogy is optical. White light contains all frequencies. A prism separates what is already present. The temporal prism does not add knowledge the model lacks; it selects which knowledge dominates.

We propose to test this with Abraham Lincoln—a figure whose documentary record is extensive, well-studied, and spans a dramatic intellectual evolution.

Three Conditions

Condition 1: The Contemporary Lens. Protocol files plus structured questions, seeded with Jon Meacham’s *And There Was Light* (2022). A modern interpretive biography that narrates Lincoln through a contemporary moral arc.

Condition 2: The Primary Source Lens. Protocol files plus the same questions, seeded with a curated corpus of Lincoln’s own writing. The selection of *which* primary sources is itself a critical methodological question. Candidates include the Lincoln-Douglas debates, the letter to Horace Greeley, the Emancipation Proclamation, the Gettysburg Address, the Second Inaugural, and selected private correspondence. But are primary sources the optimal temporal prism? Lincoln’s prose is often strategically ambiguous in ways that may not instantiate as a coherent voice. How does one build a small corpus to *hone* the larger corpus already in the model?

Condition 3: The Bare Model. Protocol files plus the same questions. No seeding material. The control condition—the unanchored instantiation most likely to produce the culturally dominant Lincoln rather than the historically disciplined one.

Methodology

Each condition is executed by an independent coding agent—an autonomous LLM instance operating within a directory containing only the protocol files, seeding materials, and structured questions. The agent convenes and interrogates. It does not improvise or bring ideological priors. This removes the convener’s influence from the instantiation, isolating the variable under investigation: the effect of the temporal prism.

An orchestrator agent relays questions and applies structured pushback. This sacrifices the pedagogical richness of live dialogue but gains reproducibility: three conditions, identical questions, controlled variables.

Dual-Lens Evaluation

Outputs are evaluated along two axes. The **historical axis**: Does the instantiated Lincoln reason in ways consistent with the documentary record? Does the persona exhibit the anticipated intellectual nuance that evolves from cautious constitutionalist to emancipator depending on the timepoint of refraction? The **computational axis**: Do the three conditions produce statistically distinguishable outputs in vocabulary, embedding similarity, and argumentative structure? The intersection—where historical judgment and computational measurement converge or diverge—is where the most important findings will emerge.



Part V: The Second Committee

The critical memo proposes a second deliberative council to address what the Committee of Patriots could not: the convergence of multiple crises—democratic erosion, wealth concentration, ecological strain, epistemic fragmentation, imperial overreach—into a single compounding conjuncture. A *polycrisis*.

The Committee of Patriots asked: *What should citizens own?* The second committee asks a prior question: **What must citizens confront?**

The proposed roster creates a deliberate collision between institutional reform and prophetic rupture. On one side: **Lincoln**, whose free labor doctrine and wartime constitutionalism tested whether democratic institutions can survive crises of their own making; **Theodore Roosevelt**, whose trust-busting nationalism attempted to discipline corporate power in the name of democratic vitality; and **Franklin Roosevelt**, who expanded executive authority enormously but in a distinctly democratic register—as stabilization, not domination.

On the other side: **William Lloyd Garrison**, who rejected incrementalism and denounced the Constitution itself, pressing the question of whether participation in a structurally unjust system legitimates it; **Eugene Victor Debs**, who insisted that political democracy without economic democracy is hollow; and **Dr. Martin Luther King Jr.**, whose late-period synthesis recognized that militarism, racial subjugation, and economic exploitation form a single interlocking system.

King's presence is essential because it ensures the deliberation cannot remain within the frame of American institutional reform. His concept of the beloved community was not merely a domestic ideal but a vision of transnational justice. The Founders built the imperial architecture. They can diagnose oligarchy but not empire. Garrison, Debs, and King see what the Founders cannot see about themselves.



Part VI: What Remains

Open Questions Requiring Collaborative Resolution

The path from the critical memo to Protocol v3.0 is visible, but several questions cannot be resolved from within the project alone. They require the kind of deep historical judgment that prompted this exchange in the first place. Here is a summary of the open questions and the input needed to proceed.

1. The Lincoln Prism: Which Sources, Which Moment?

The Prism Experiment's primary source condition requires a curated corpus. But which Lincoln are we trying to evoke? The prairie lawyer of 1858? The wartime president of 1863? The theological voice of the Second Inaugural? The selection of sources is itself an interpretive act—and it determines what kind of Lincoln the model will instantiate. We need guidance on which documentary materials constitute the most revealing temporal lens, and whether primary sources alone are sufficient or whether a secondary interpretive frame (a specific biography, a historiographical essay) produces a more coherent vessel.

2. Questions for the Lincoln Interrogation

The structured questions posed across all three conditions must be designed to reveal differences in reasoning, not just differences in style. We need questions that a historian would recognize as diagnostic—questions where Lincoln's answer would differ meaningfully depending on which moment in his evolution is instantiated. What questions would most sharply differentiate the cautious constitutionalist from the emancipator, the free-labor pragmatist from the Second Inaugural theologian?

3. The Second Committee: Central Questions for the Polycrisis

The roster is compelling. But a committee needs a charge. The Committee of Patriots was given a specific question: what should citizens do to reclaim economic sovereignty? The second committee needs an equally precise framing. What are the two or three central questions that would most productively generate collision between the institutional reformers and the prophetic voices? How do we frame the charge so that

the tension between reform and rupture becomes generative rather than merely oppositional?

4. Evaluating Historical Fidelity

The Prism Experiment produces three sets of Lincoln’s responses. The computational evaluation is straightforward: NLP methods can measure vocabulary divergence, embedding distances, rhetorical structure. But the historical evaluation is the one that matters most. What criteria would a historian use to assess whether an instantiated Lincoln “tracks with the documentary record”? What would constitute evidence of authentic historical reasoning? Is there a methodology we should formalize here—a rubric for evaluating persona fidelity from the discipline of history?

5. The Counter-Voice Architecture

The recommendation to embed historiographical plurality into the protocol is powerful. But what does this look like in practice? Is the counter-voice a separate summoned persona (Frederick Douglass paired with Jefferson)? Is it a structured adversarial role written into the protocol specification? Is it a set of prompts that force the convener to surface blind spots before the deliberation begins? There are several possible architectures here, and the right one depends on what historians would consider methodologically sound.



Something obvious has become visible in this exchange—something Silicon Valley is desperately missing. The language models contain the raw material of centuries of human thought. But the people building and deploying them are not, by and large, the people trained to read, contextualize, and interrogate text with the discipline that the humanities demand. What has happened here—a historian’s methods of reading applied to AI infrastructure—should not be rare. It should be the norm.

The Dossier was released as an invitation. The Halpern memo accepted it. The conversation has begun, and what it produces next depends on whether the collision continues.

The Word against the Flood. The threshold widens.



Jacob E. Thomas, PhD
Austin, Texas • February 2026

